

# DNS Data Exfiltration Prevention: Kernel-Enforced Endpoint Security

*Scalable Framework to Disrupt  
DNS C2 and Tunneling*

Vedang Parasnis

University of Washington

# Agenda






- Data Exfiltration / Data Breaches Phases
- DNS attack vectors to exfiltrate data
- Drawbacks of current approaches
- Security Framework Overview
- Kernel Enforced Endpoint Security Architecture (Kernel + Userspace)
- Results, Evaluation, Demo
- Protect Linux Kernel from malicious tampered endpoint security eBPF programs
  - Cloud CA + kernel keyring + BPF LSM + Kernel Datapath

# Data Exfiltration / Data Breaches


**Definition:** Unauthorized extraction or transmission of sensitive data from a system

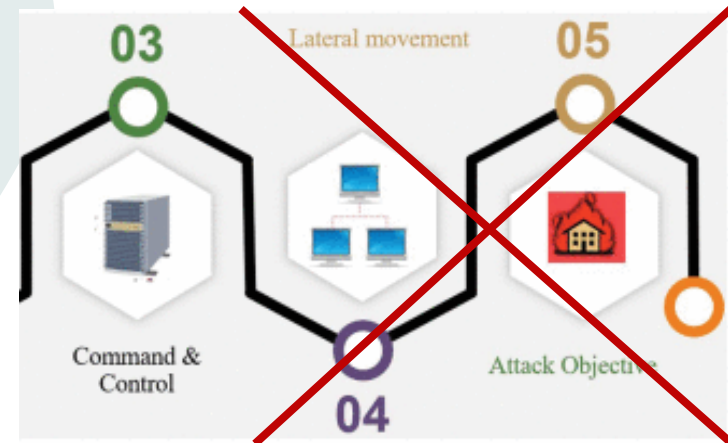
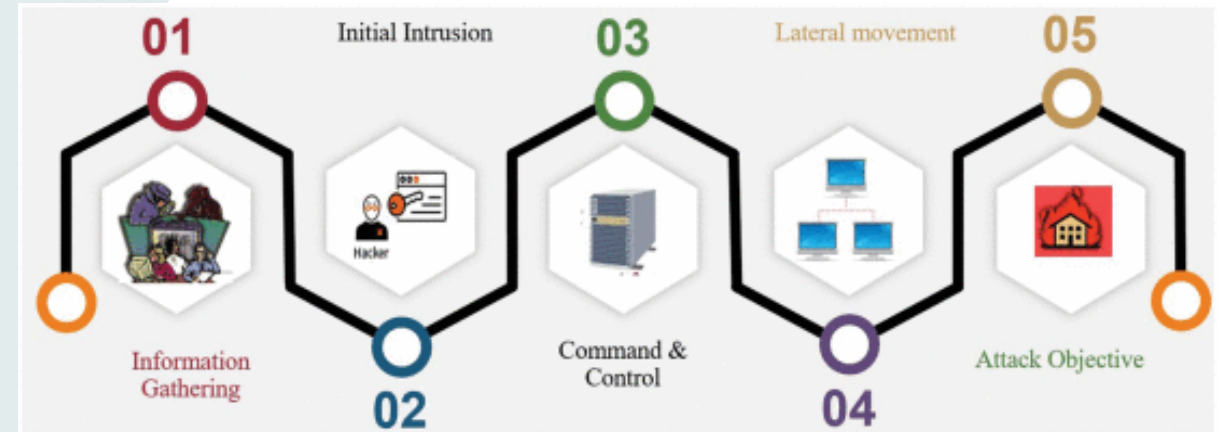
**Impact:** Reputation, Financial damages

## Attack Lifecycle

-  Information Reconnaissance
-  Initial Intrusion / Infiltration
-  **Command and Control**
-  Lateral Movement
-  Command Execution and Data Breaches

## Core Defense Strategy

-  Endpoint Security (EDR / XDR)



# DNS Data Exfiltration

**DNS C2** - Uses DNS queries and responses to maintain covert communication with attacker infrastructure.

**DNS Tunneling** - Encapsulates arbitrary data within DNS packets to bypass network restrictions.

**DNS Raw Exfiltration** - Leaks sensitive data files directly in DNS queries.



Malware sends username and password data encoded in base64 as hostname label

- **Remote Code Execution (RCE)**
  - Shell code exploits
  - Script executions, File corruptions
  - Process Side channeling exploits
  - Example: **Sliver C2, Hexane, APT29 (Cozy Bear), Skitnet.**
- **Persistent Backdoors**
  - Deployment rootkits, ransomwares
  - Example: **Turla group**
- **Network Pivoting (Port Forwarding)**
  - Compromised machines act as proxies to reach deeper into private infrastructure
  - Example: **Cobalt Strike, Hexane, DNSSystem**

# Existing Approaches

- **Semi-Passive Analysis**
  - DNS Exfiltration Security as Middleware
- **Passive Analysis**
  - Anomaly Detection
  - Threat Signatures, Domain Reputation scoring

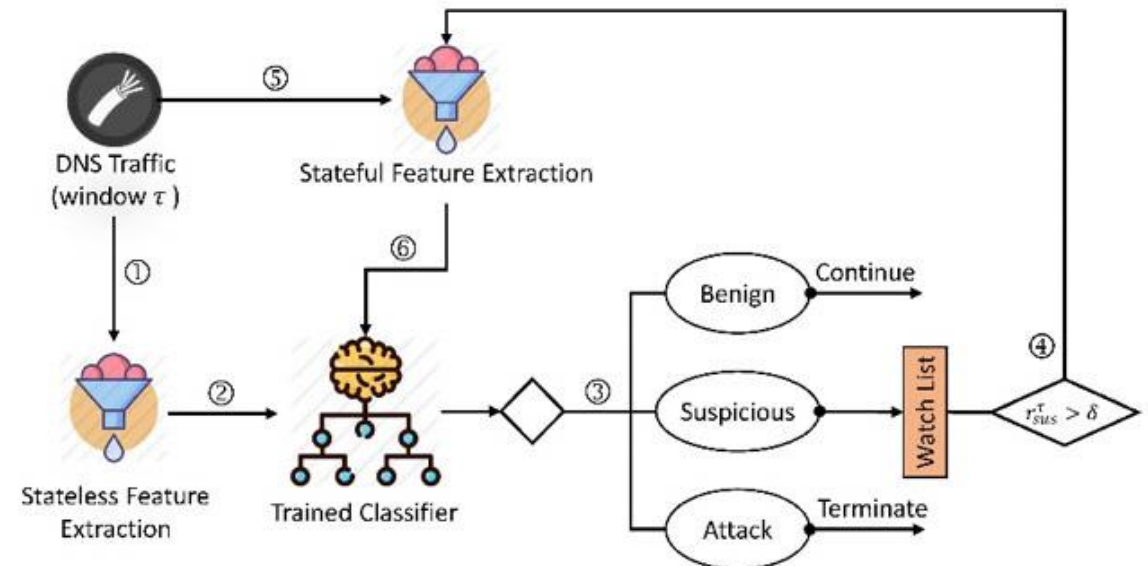


# Existing Approaches – Passive Analysis

- **Anomaly Detection:**
  - **Traffic Behavior Analysis**
    - DNS Passive Traffic Volume Analysis
    - DNS Passive Traffic timing Statistical Analysis
  - **Machine Learning-based Threat Intelligence**
    - Uses machine learning models to identify traffic anomalies.
- **Threat Signatures:**
  - DNS Domain Scoring
  - Malicious domain signature

**Stateless Features** – Lexical Analysis

**Stateful Features** – Statistical Analysis



# Issues with current approaches

- **Slow Detection → High Dwell Time → More Damage**
- **Extremely slow to Advanced C2 Attacks**
  - More Damage if C2 infrastructure employs multiplayer mode (Botnet of C2 server exploiting scaled environments)
- **Dynamic Threat Patterns:**
  - Varying Throughput
  - Slow and Stealthy Rate
  - Kernel Encapsulated Traffic
  - Port Obfuscation
- **Centralized monitoring and analysis systems don't scale**
- **Ineffective over IP Masquerading & Domain Generation Algorithms**

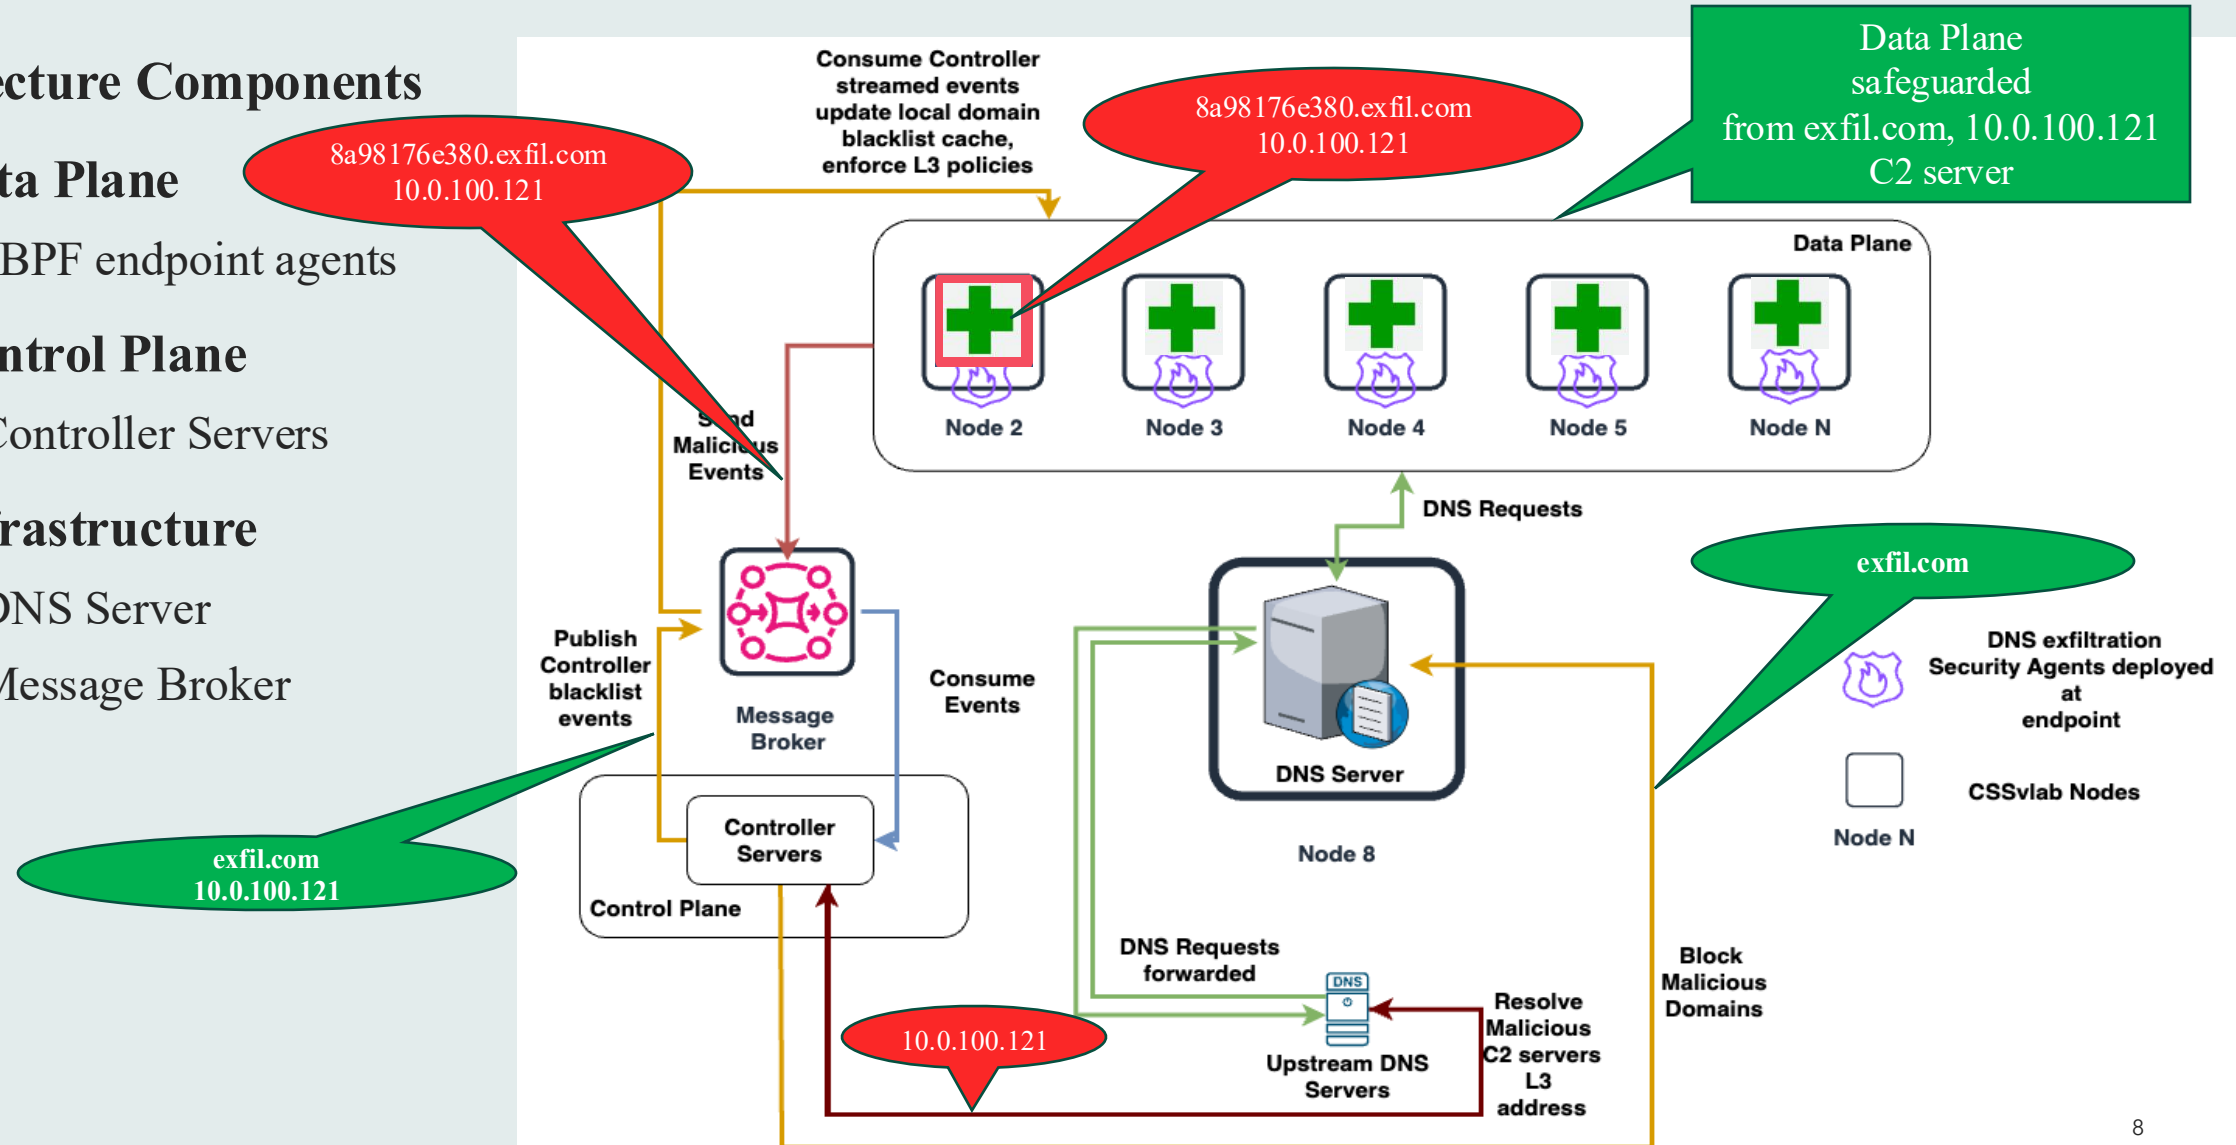
**Solution:**

**Real-time, proactive enforcement at Ring 0 — inside the kernel.**

# DNS Exfiltration Security Framework Overview

## Architecture Components

- **Data Plane**
  - eBPF endpoint agents
- **Control Plane**
  - Controller Servers
- **Infrastructure**
  - DNS Server
  - Message Broker





# Security Framework Goals

## **Disrupt DNS covert C2 channel attacks, data exfiltration.**

Implement in-kernel deep packet inspection and enforcement to block all forms of DNS exfiltration channels.

## **AI-Assisted Threat Detection**

Use deep learning in userspace to detect advanced obfuscated exfiltration payloads with high accuracy aiding kernel network enforcements.

## **Malicious Process Aware Active Response (Threat-Hunt and Kill)**

Link exfiltration attempt to parent process and kill implants processes, preventing lateral movement and further damage.

## **Dynamic Cross-Layer Policy Enforcement**

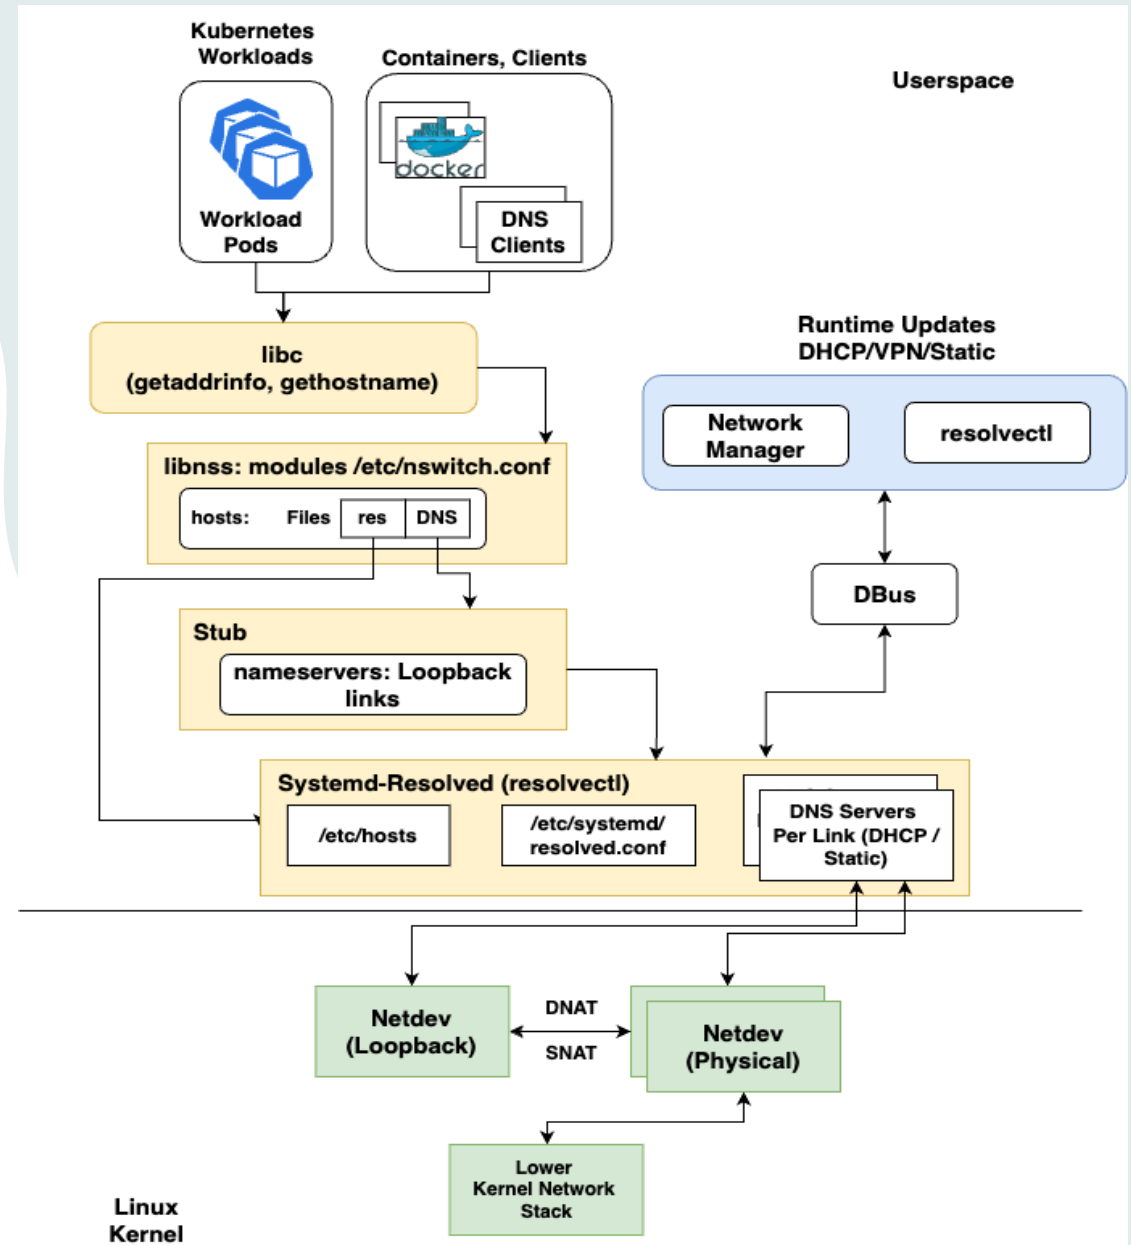
Enforce in-kernel L3 network policies adaptively and domain blacklisting on DNS server to combat DGA.

## **Scalable Multi-Cloud Deployment**

Ensure framework's horizontal scales for real-world production cloud environments.

# Systemd-Resolved

- **Userspace**
  - Libc (dns\_utils)
  - Libnss (nss modules (nss-dns, nss-myhostname))
  - Systemd-resolved (resolvectl)
  - System Daemons
    - Network Manager (DHCP)
    - Dbus
- **Kernel**
  - Network Stack each netdev (east-west, north-south traffic)



# Kernel Enforced Endpoint Security

## Agent based Endpoint Security

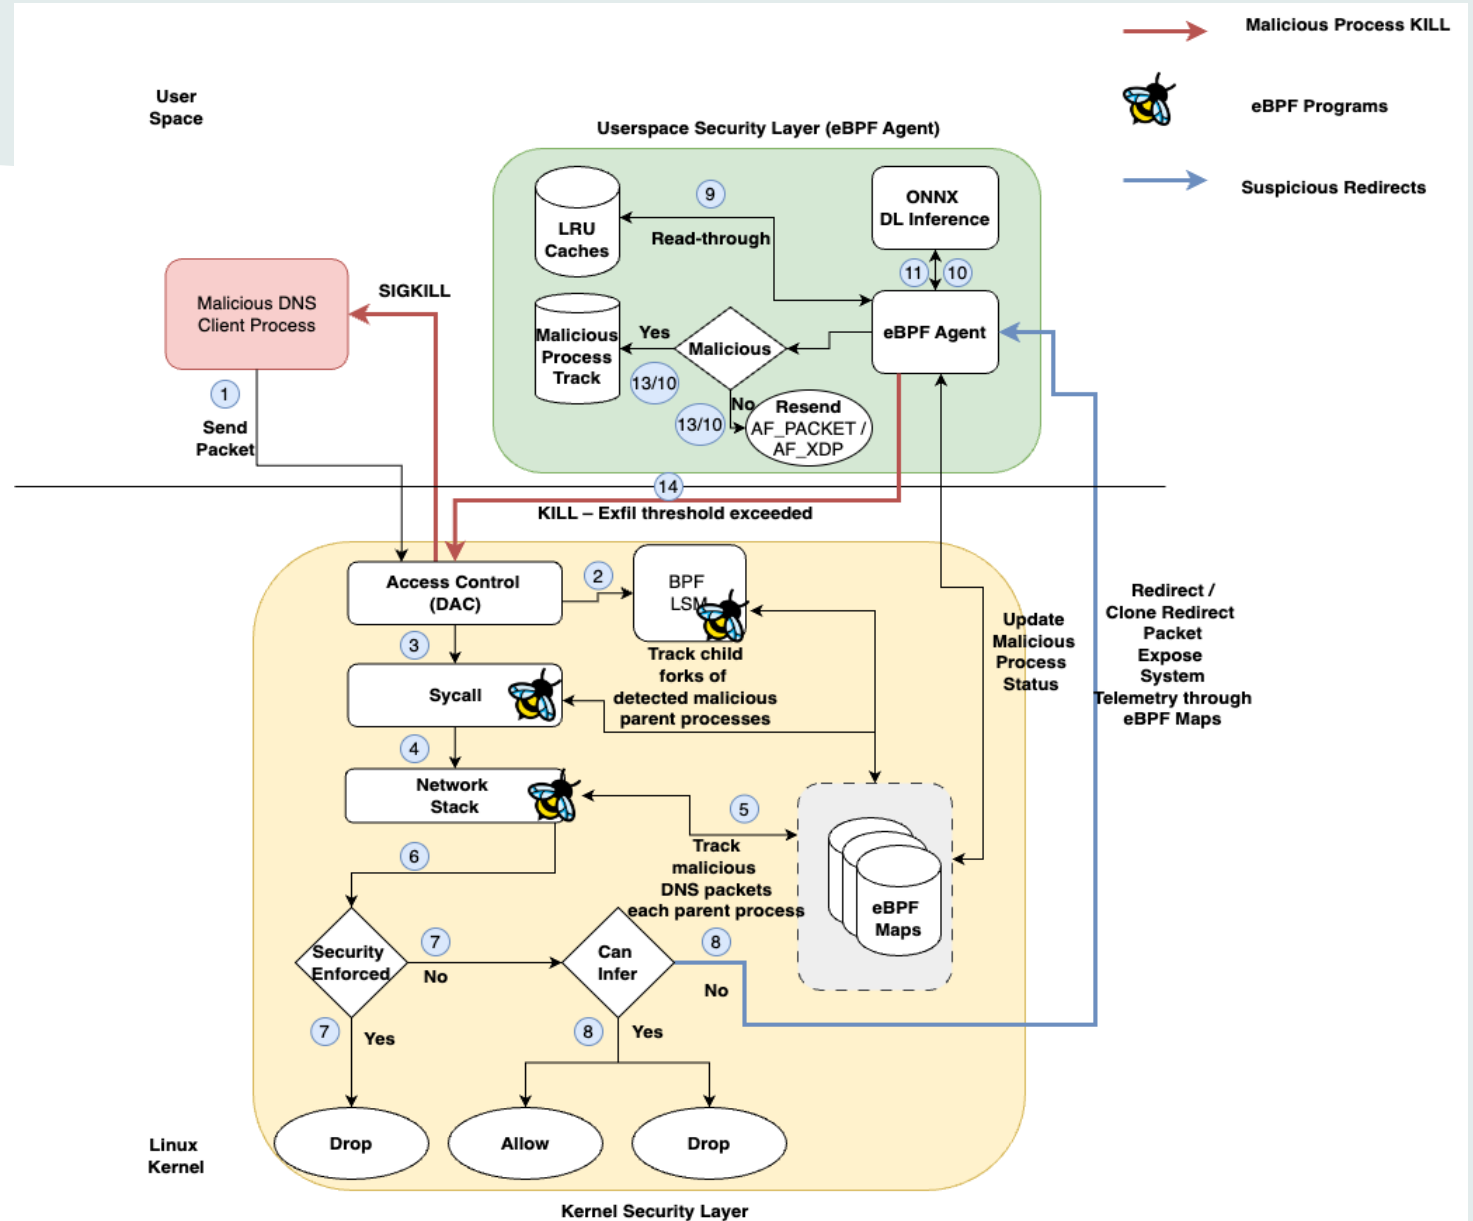
### Continuous Security Enforcement Event Loop

#### Userspace

- eBPF Agent
- eBPF Agent LRU Caches
- ONNX Quantized Deep Learning Model
- Kernel malicious metrics export

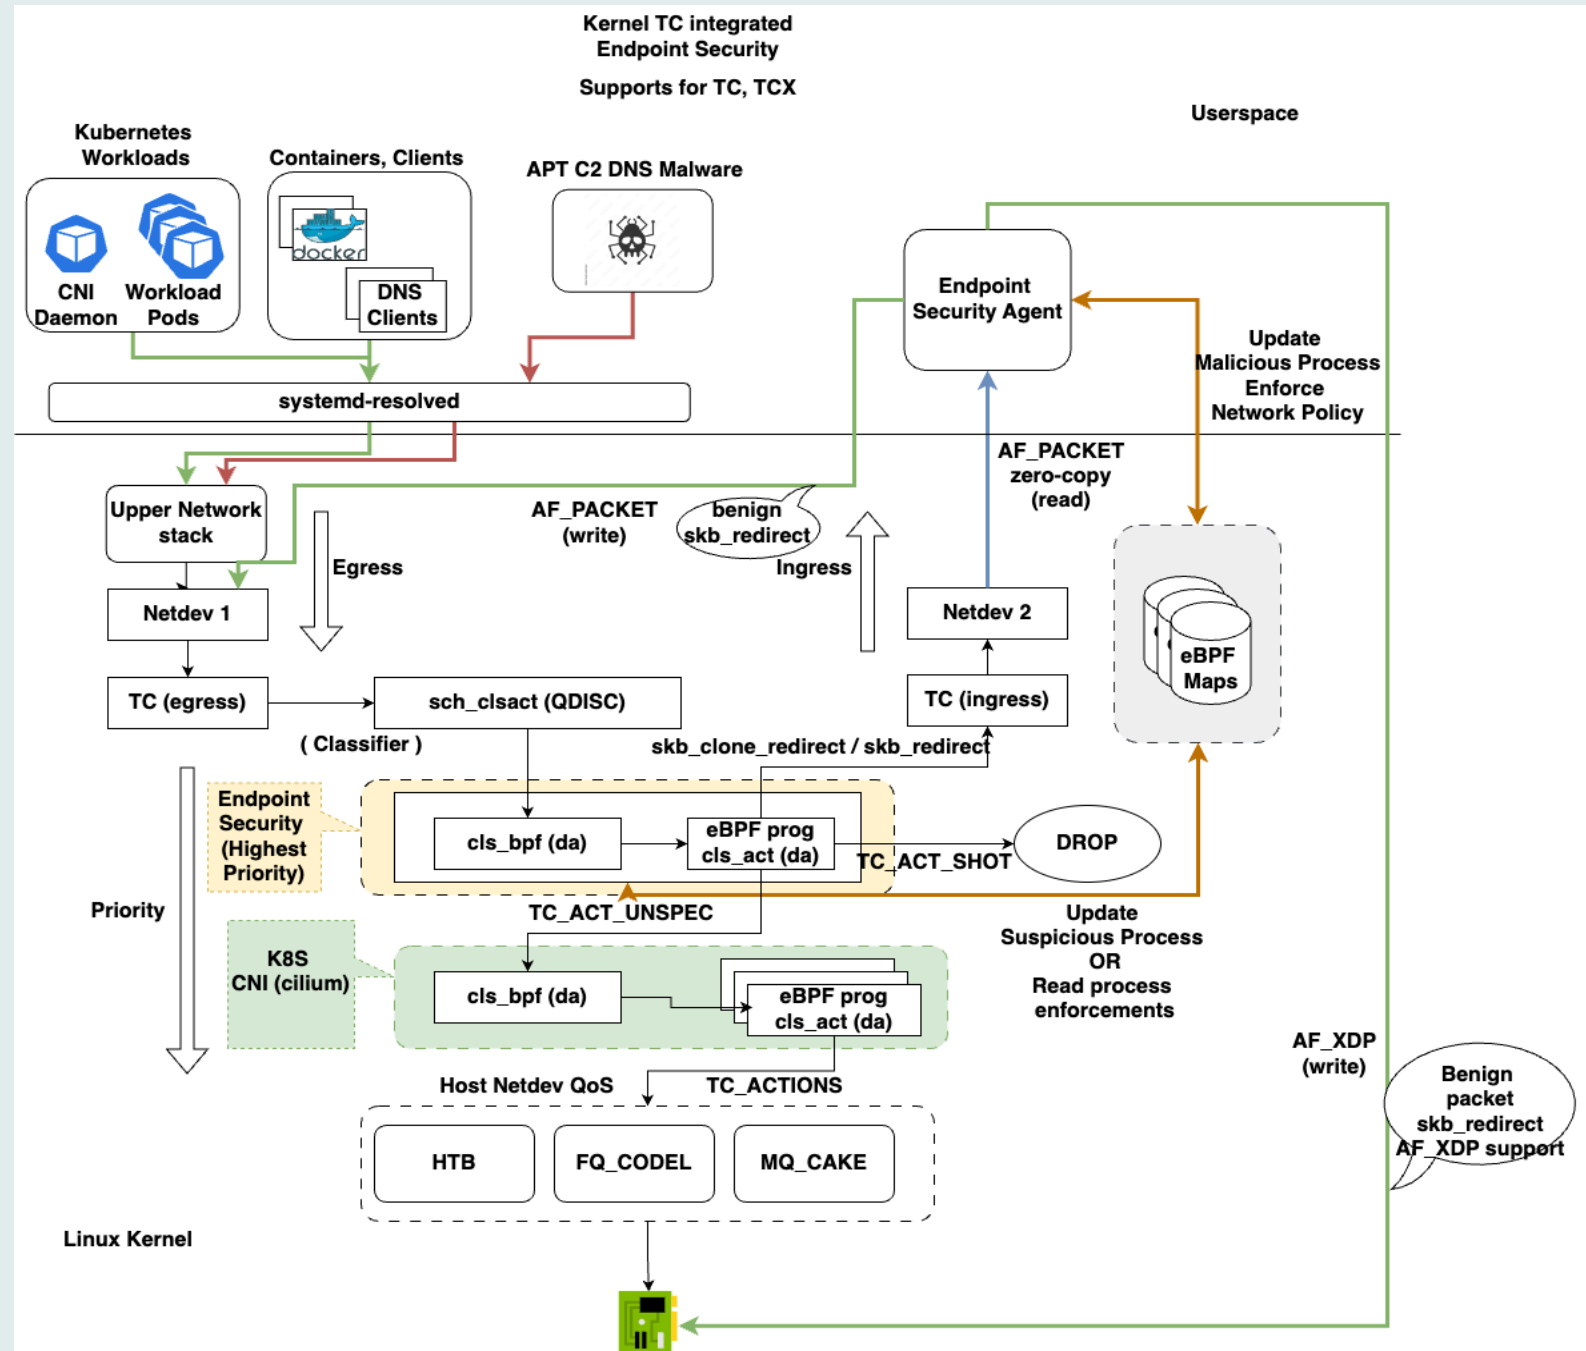
#### Linux Kernel

- Inference Unix Domain Sockets
- eBPF Ring Buffers (malicious events)
- Network Stack (eBPF programs)
  - Socket Layer
  - Traffic Control
- Access Control Layer (eBPF programs)
  - Security Modules (eBPF LSM)
  - Syscall (eBPF Tracepoints)



# Kernel Datapath Enforcement Layer

- Sockets
- TCP/IP Stack
- Netfilter
- Traffic Control (QoS)
- Network Drivers (XDP)



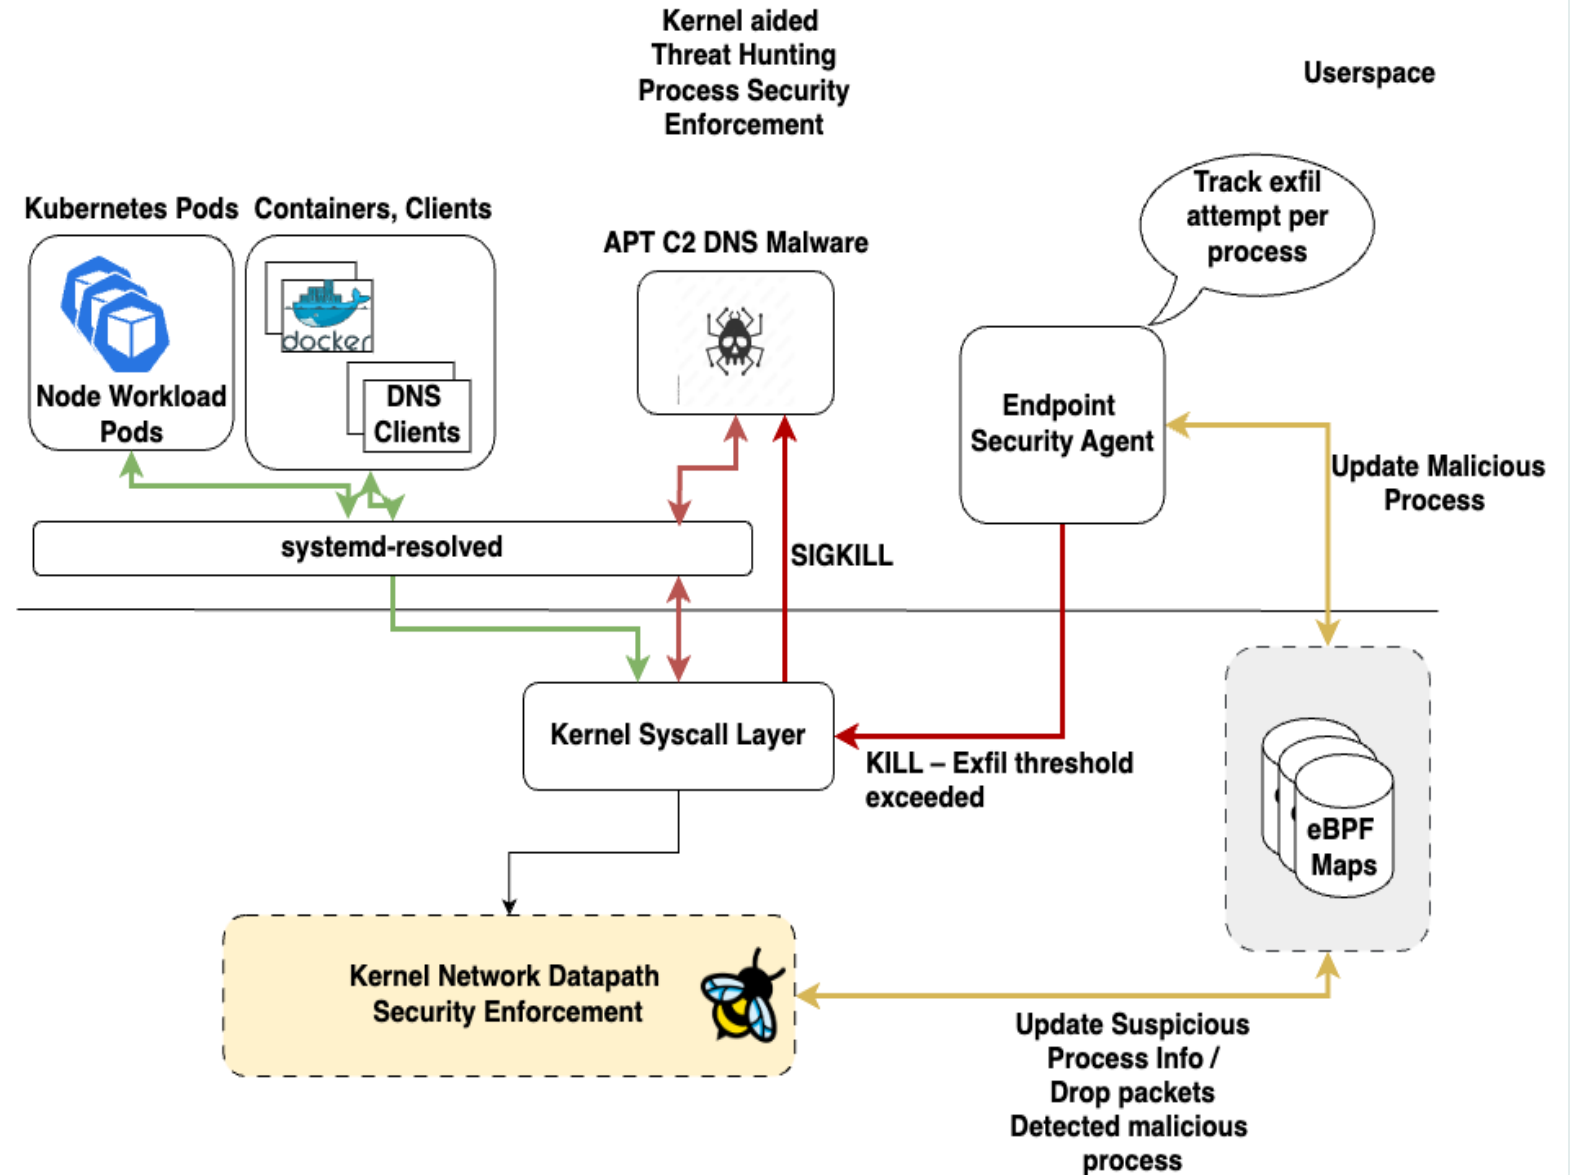
# Process Enforcement Layer

## Userspace

- Send SIGKILL to malicious process

## Linux Kernel

- Kills the malicious implant instructed by userspace endpoint security agent.



Linux Kernel

# eBPF Endpoint Agent Operations Modes

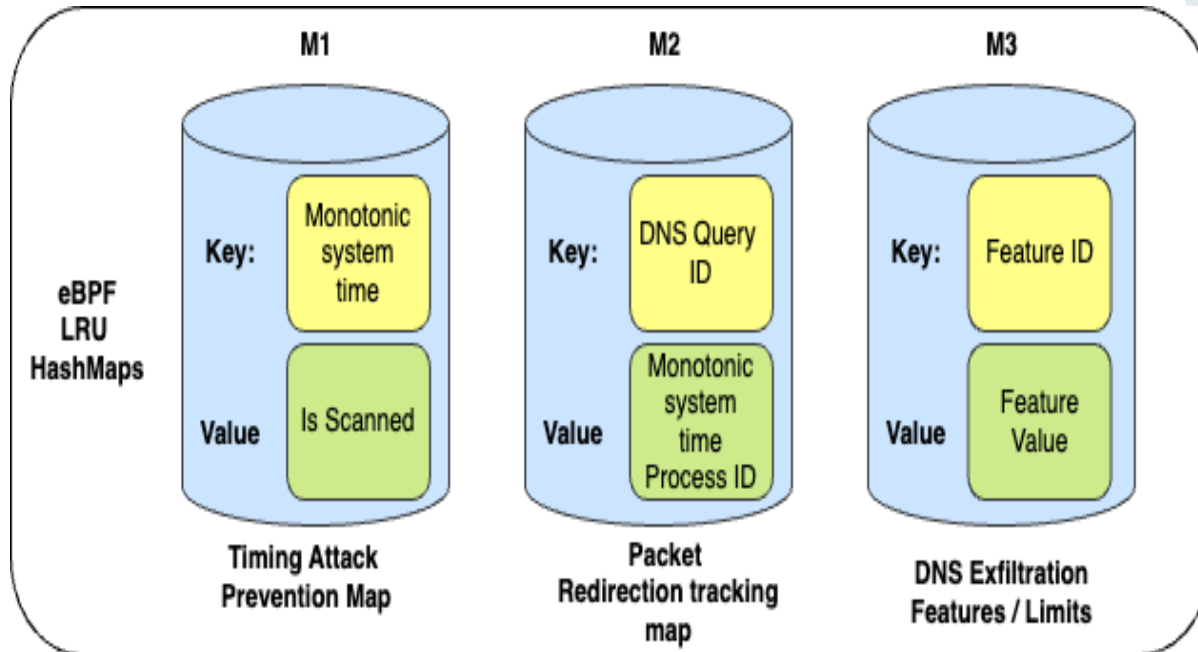
eBPF Agents in Data Plane handle DNS exfiltration over UDP

Mode	Goal	Requirement	Security Enforcement Process
Strict Enforcement Active Mode	Kill C2 Implants, ensure zero data loss and C2 command execution.	DNS Traffic over UDP ports (53, 5353, 5355), for encapsulated and non-encapsulated traffic.	<ul style="list-style-type: none"><li>• <b>Kernel:</b> Live Redirects suspicious DNS packets to userspace.</li><li>• <b>Userspace</b> Trace malicious process exfiltration count and terminates it, resend benign packets.</li></ul>
Process-Aware Passive Threat Hunt Mode	Kill C2 Implants, ensure negligible data loss and minimal C2 command execution.	DNS Traffic over random UDP ports.	<ul style="list-style-type: none"><li>• <b>Kernel:</b> Allow suspicious traffic passthrough. In Kernel start threat hunting process tied to malicious DNS packets.</li><li>• <b>Userspace:</b> Trace malicious process and terminates it.</li></ul>

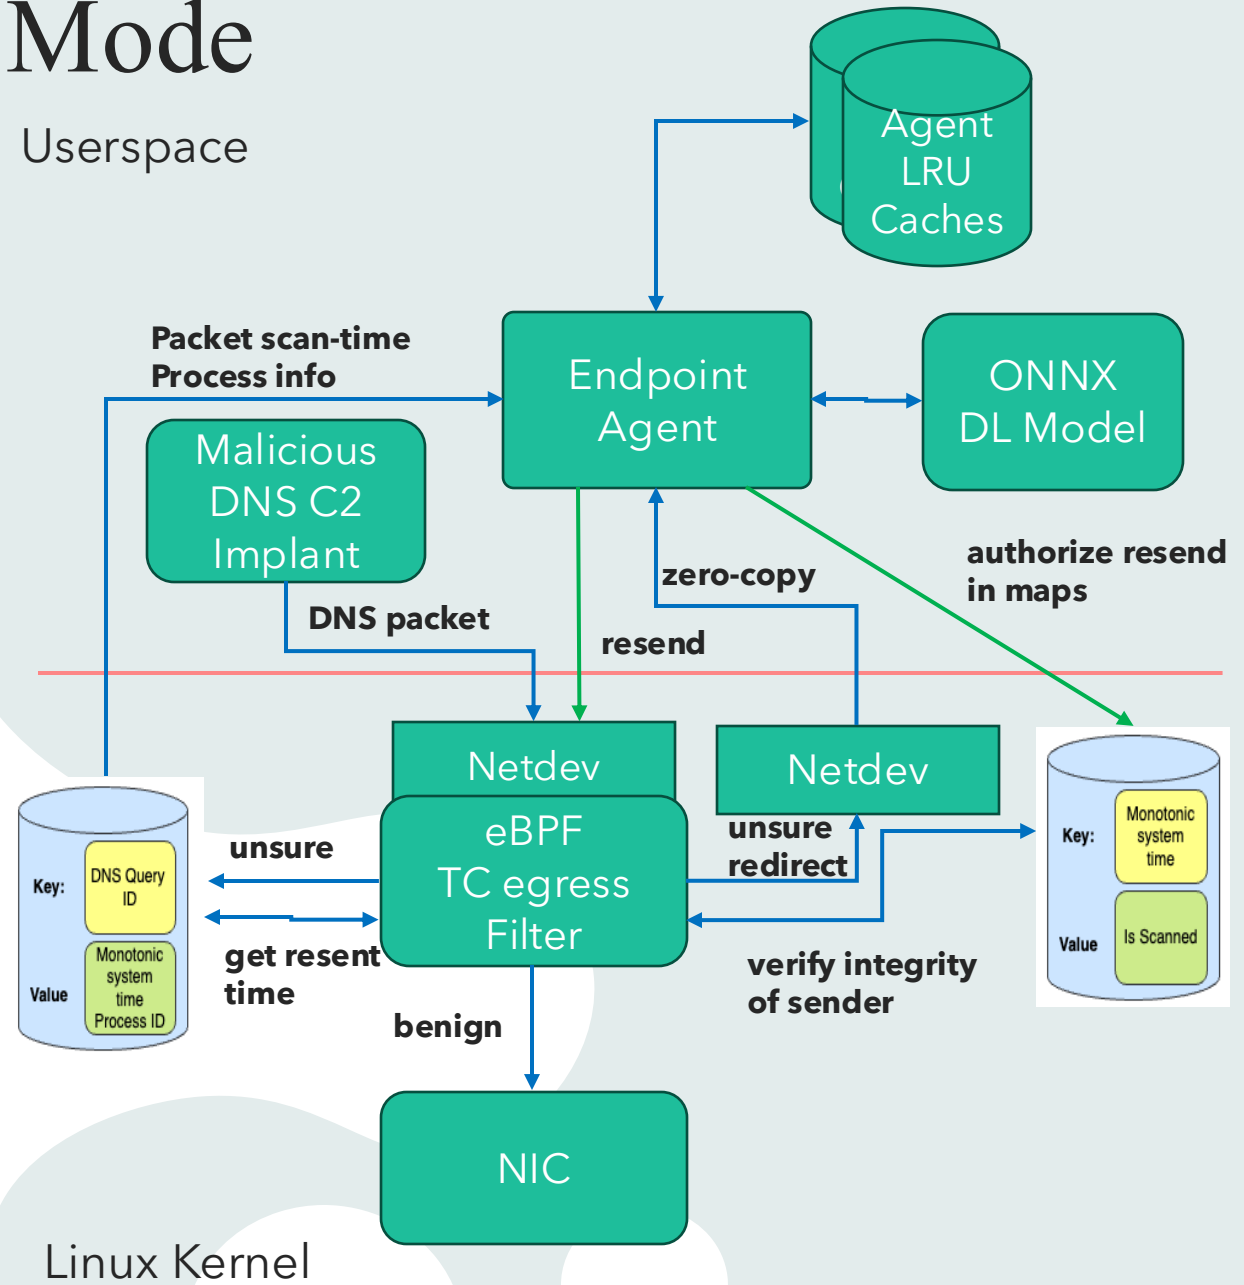


# Strict Enforcement Active Mode

- eBPF program deep parse of suspicious DNS packets from SKB
- Real-time verdict from kernel DPI eBPF program
- Userspace DL model aids classification
- eBPF timing-brute-force attack checks on resend
- Per-process exfil attempts tracked in userspace



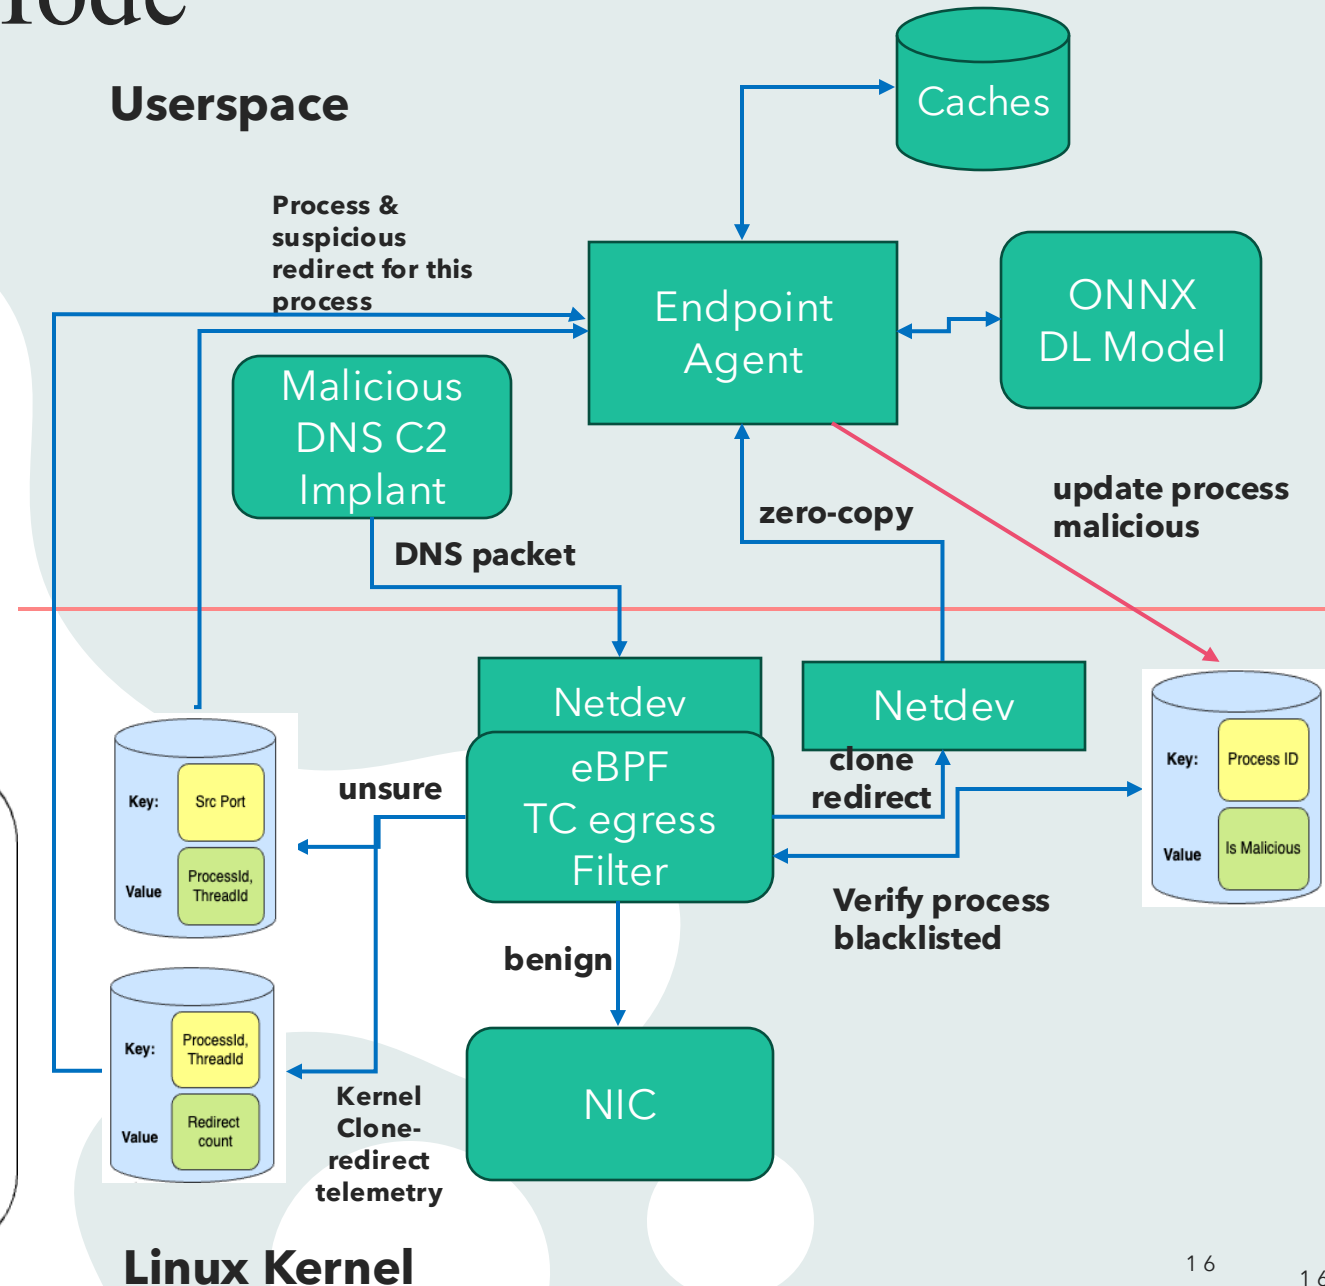
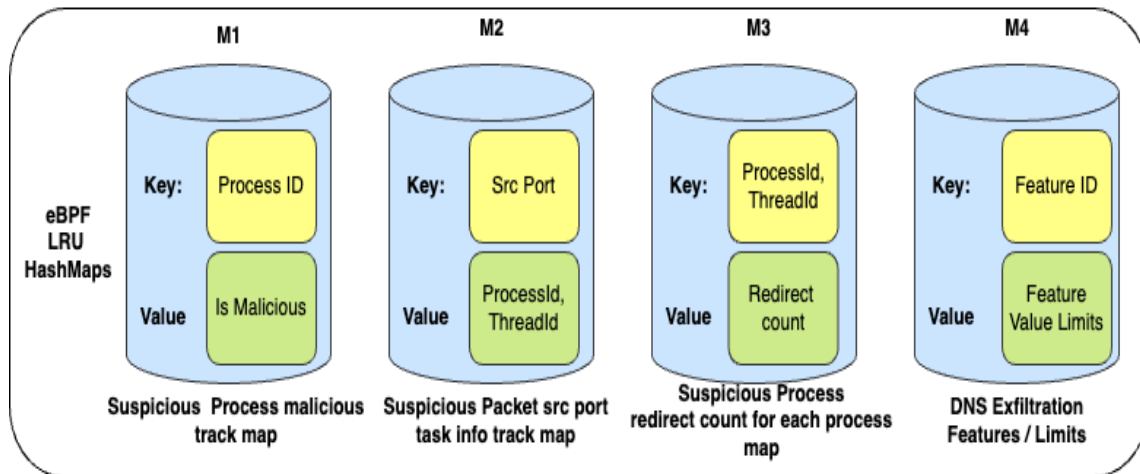
Userspace



Linux Kernel

# Process-Aware Passive Mode

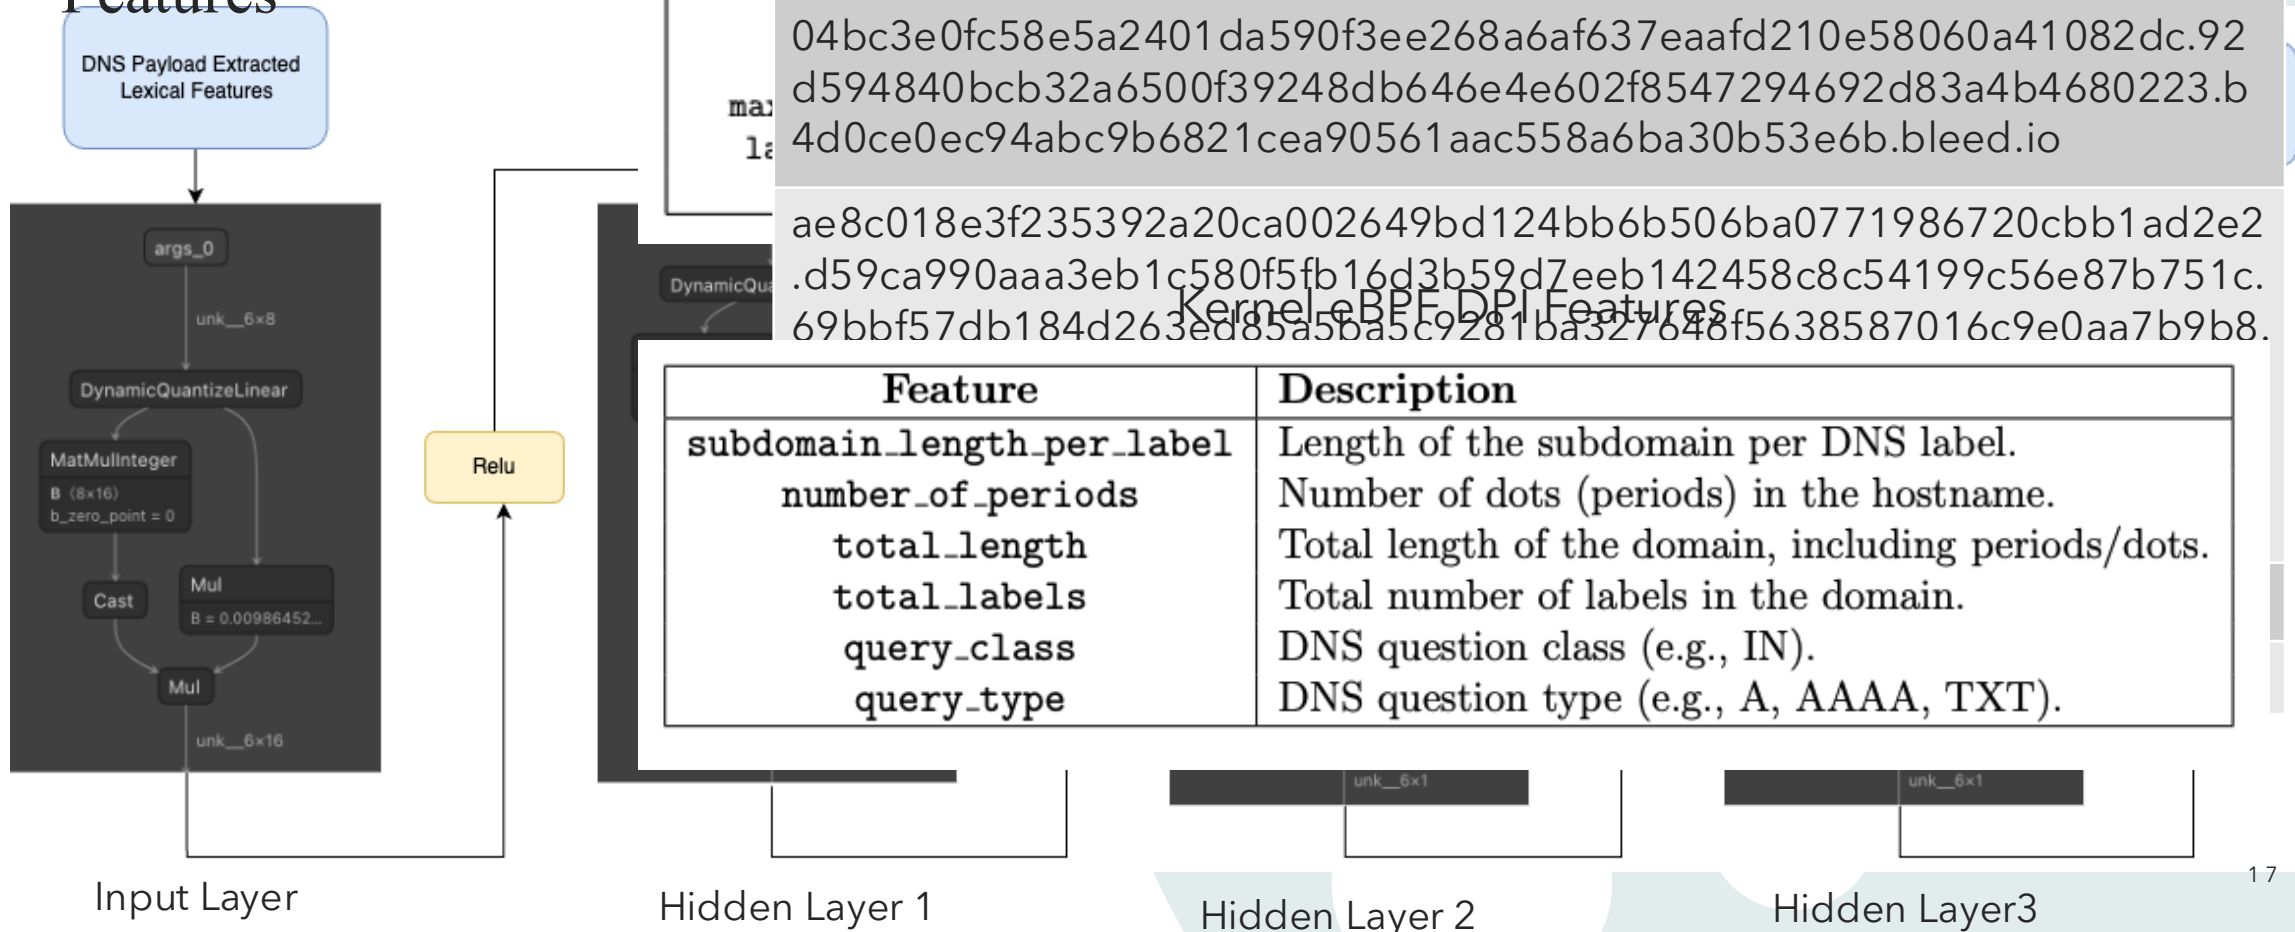
- eBPF kernel program deep parses DNS from skb
- Suspicious packets cloned to userspace
- DL model classifies and tags process
- Malicious process flagged and eBPF maps updated
- eBPF kernel program drop packet from malicious process + clone for exfiltration attempt telemetry



# DNN based DNS Data Obfuscation Detection

## DNN Features

- Model Architecture ONNX
- Sample Malicious Data
- Input Features



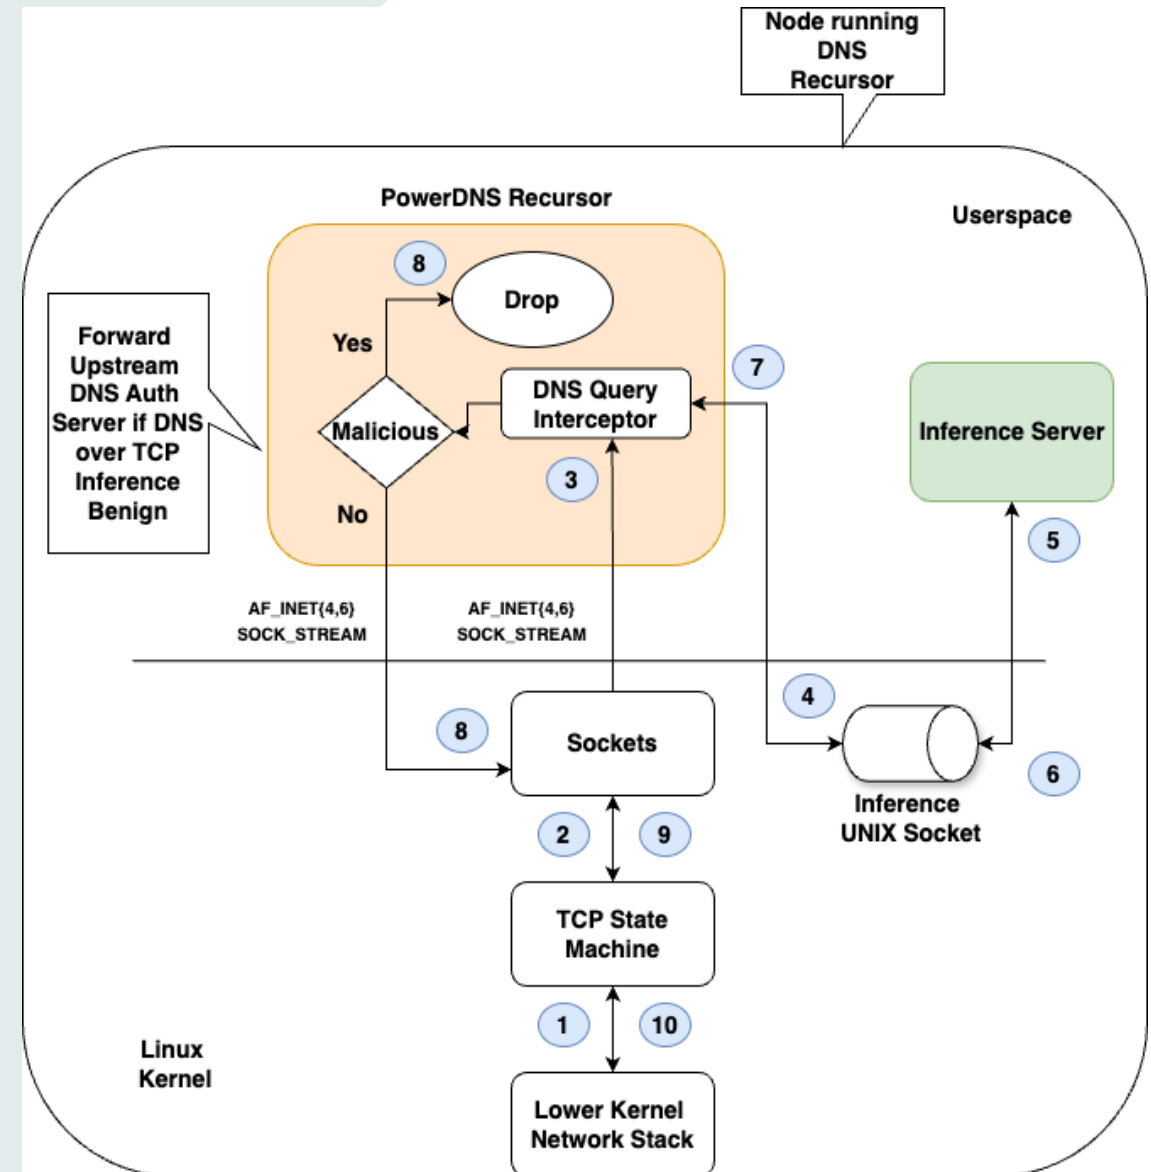
# Datasets

Dataset Type	Source / Characteristics	Size	Primary Goal
Trusted Benign Cache	Top 1M Cisco Second-Level Domains (SLDs)	1 Million	Reduce inference on known-good traffic.
ISP-Captured DNS	Live-sniffed ISP DNS traffic [ <a href="#">Ziza et al.</a> ]	50 Million	Provide real-world benign & malicious baseline.
Synthetic Exfiltration	Custom-generated (DET, DNSCat2, Sliver, Nuages, Custom Scripts, etc.);	2.4 Million	Malicious samples use varied obfuscation across file formats
Final Combined Dataset	Synthetically formed	3.8 millions	Balanced dataset w/ obfuscated payloads across file formats

# DNS Exfiltration over TCP

## Prevent DNS Exfiltration over TCP

- Runs on
  - PowerDNS Recursor
- Relies on
  - PowerDNS recursor Query Interceptors
  - Inference UNIX domain sockets



# Results and Evaluation

- Model Metrics
- Throughput comparisons (Active mode)
- Response Time per Exfiltration attempt
- Kernel DPI time (raw parse DNS protocol from SKB)
- Resources
  - Memory Usage
    - Security Agent memory usage at endpoints in data plane
  - Endpoint Agent Flame Graph

## Test Bench

CPU: Intel Xeon 6130

Memory: 8 GB

Linux Kernel: 6.12.4

Network Driver: netvsc

Bandwidth: 100 Gb/sec

Root QDISC: FQ\_Codel

Queues: 8 RX / TX

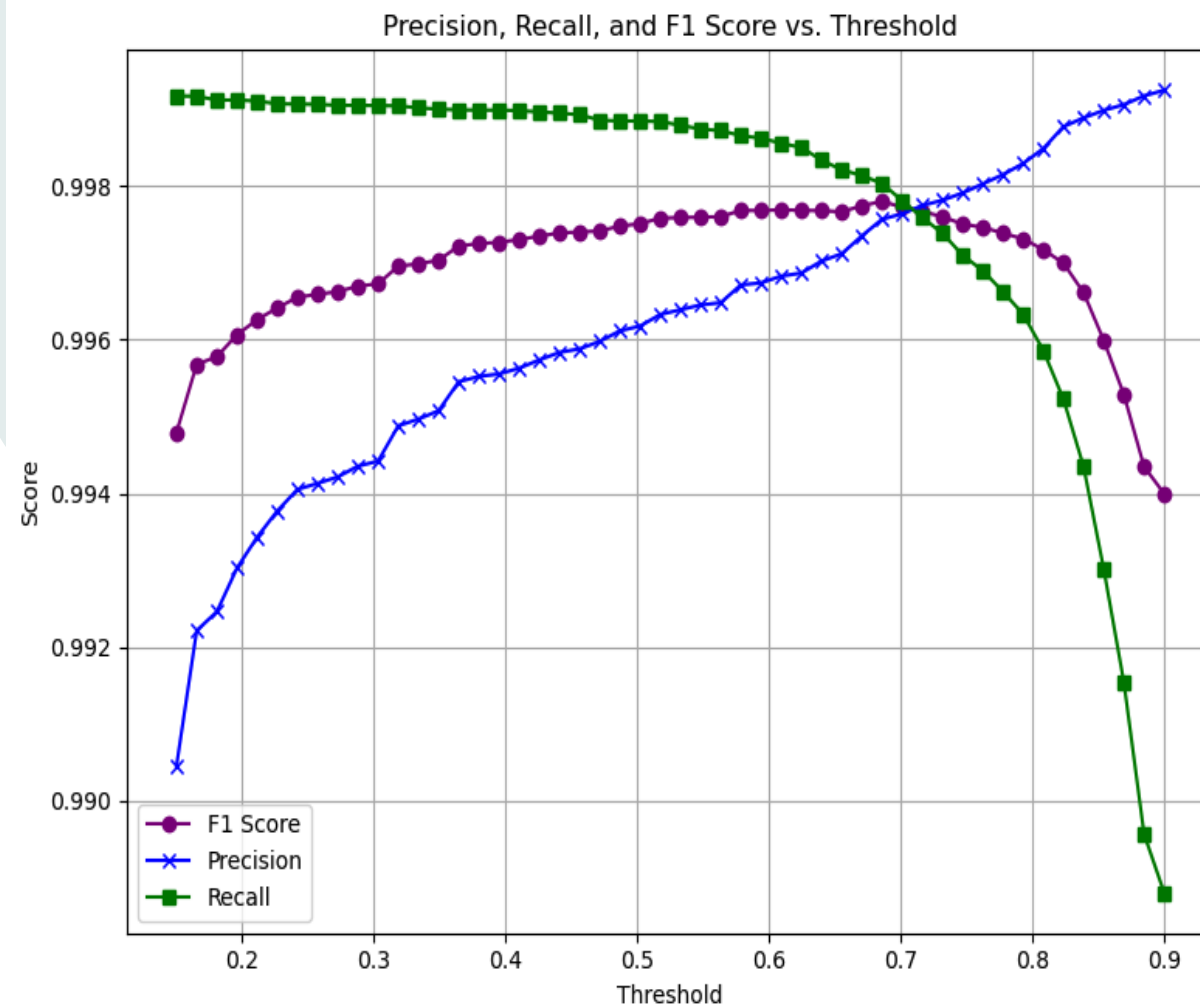


# DNN Model Metrics

Metric	Training	Validation
Accuracy	0.9973	0.9997
AUC	0.9997	0.9997
Loss	0.0099	0.0091
Precision	0.9959	0.9959
Recall	0.9987	0.9988

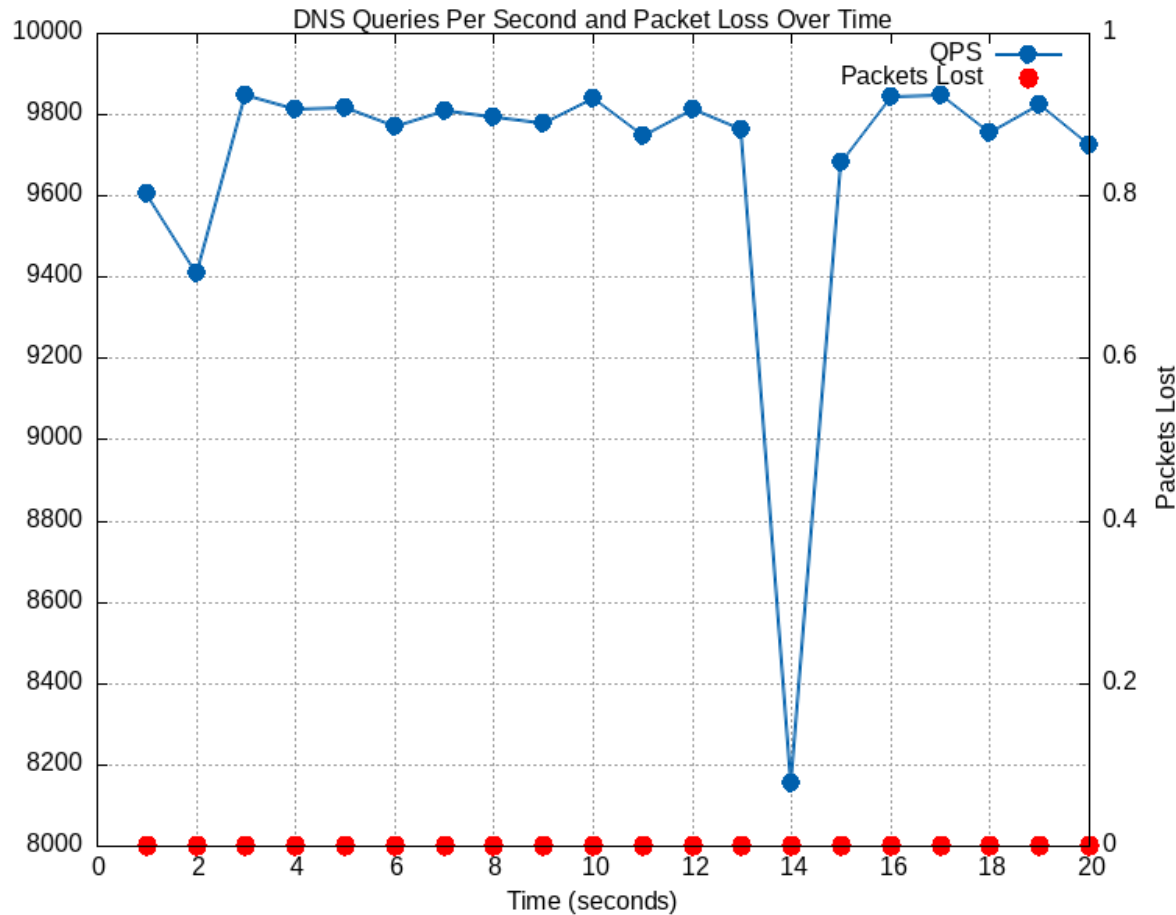
Table 5.1: Model Evaluation Metrics

Model Performance




Model Scores

# Throughput comparisons – Active Mode




Agent LRU Cache Read-Through Hit 10k DNS req/sec



ONNX Live Inferencing 10k DNS req/sec

# Throughput comparisons – Active Mode (continued)



Agent LRU Cache Read-Through Hit

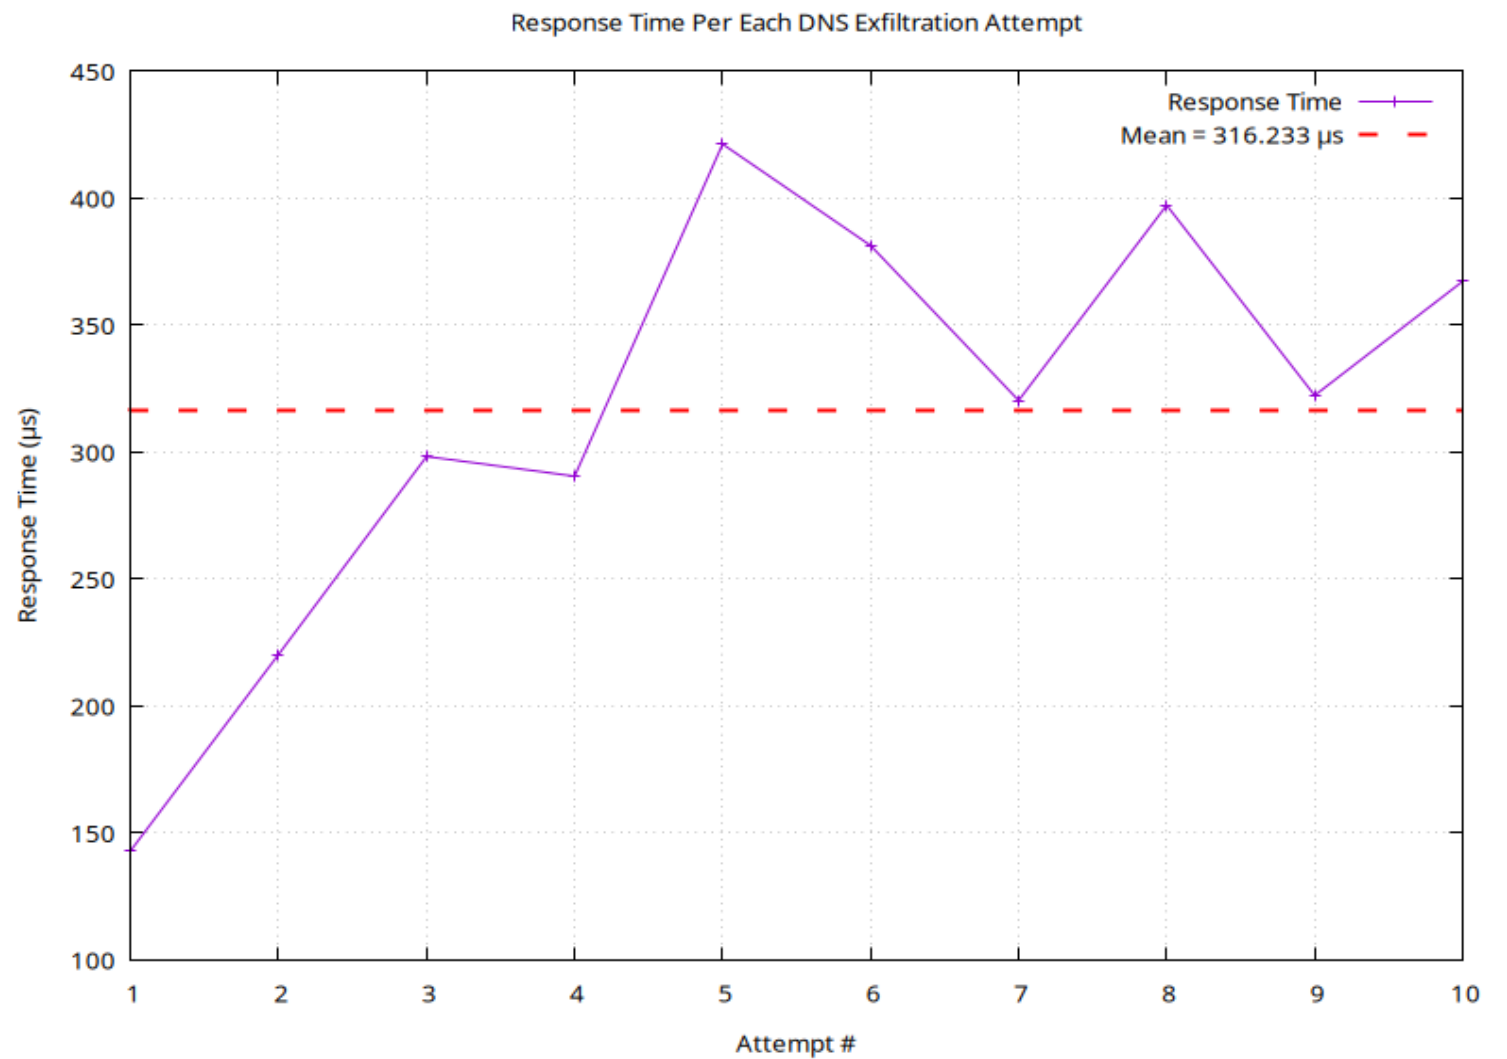


ONNX Live Inferencing

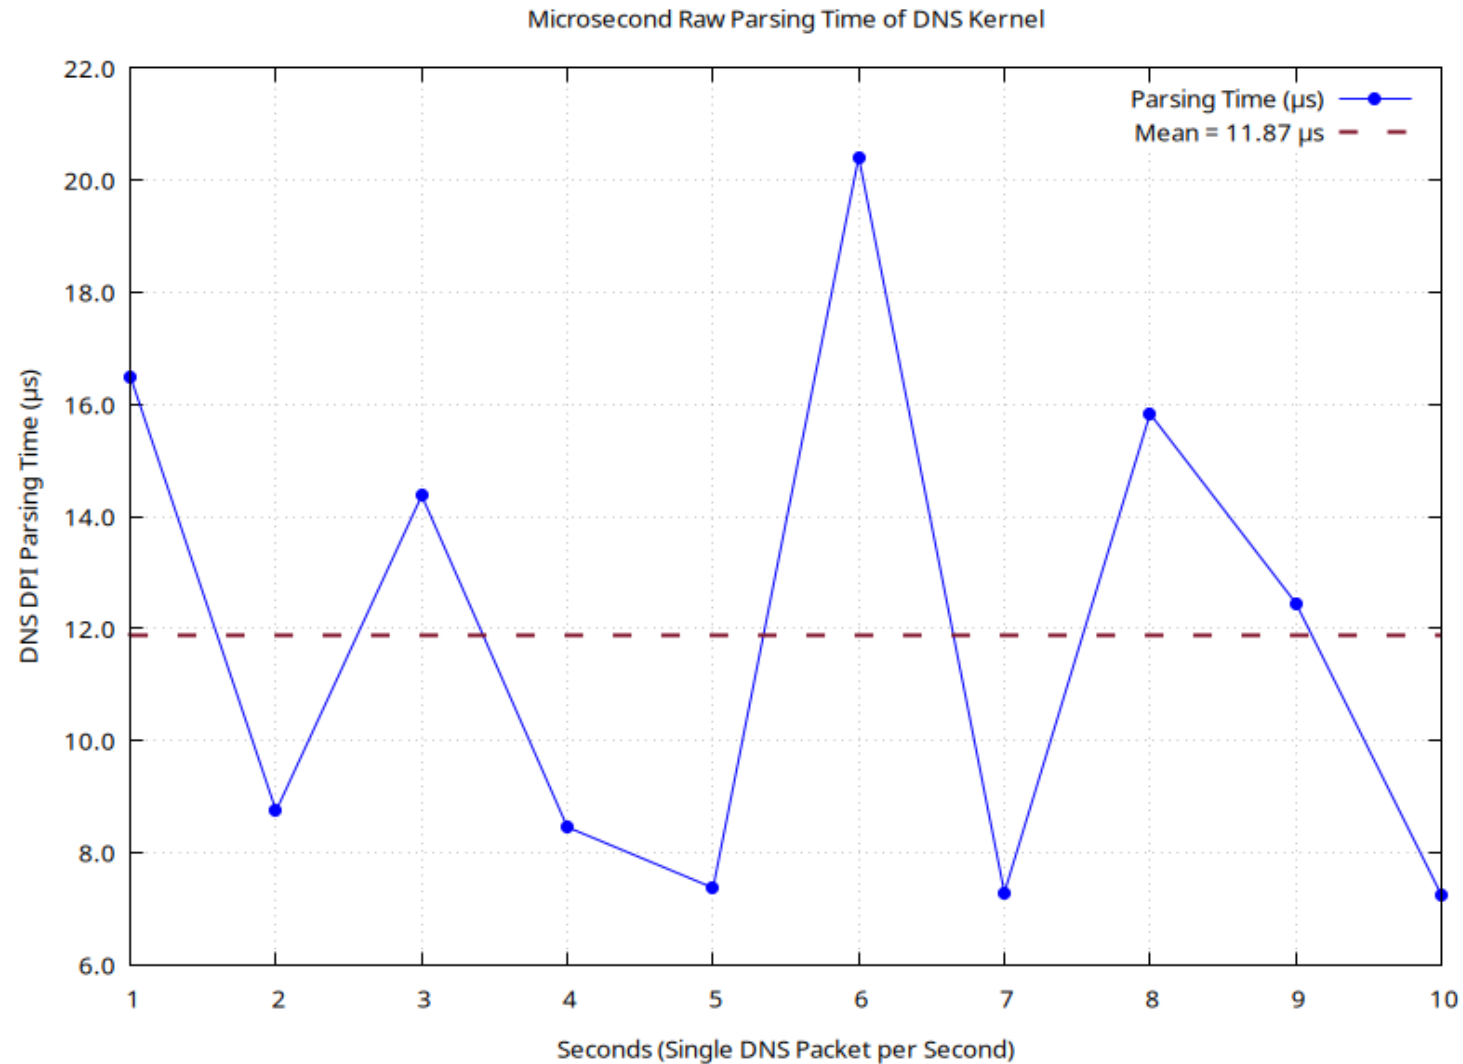
# Response Speed - Active Mode

Response Speed Before  
Implant Eventually Killed


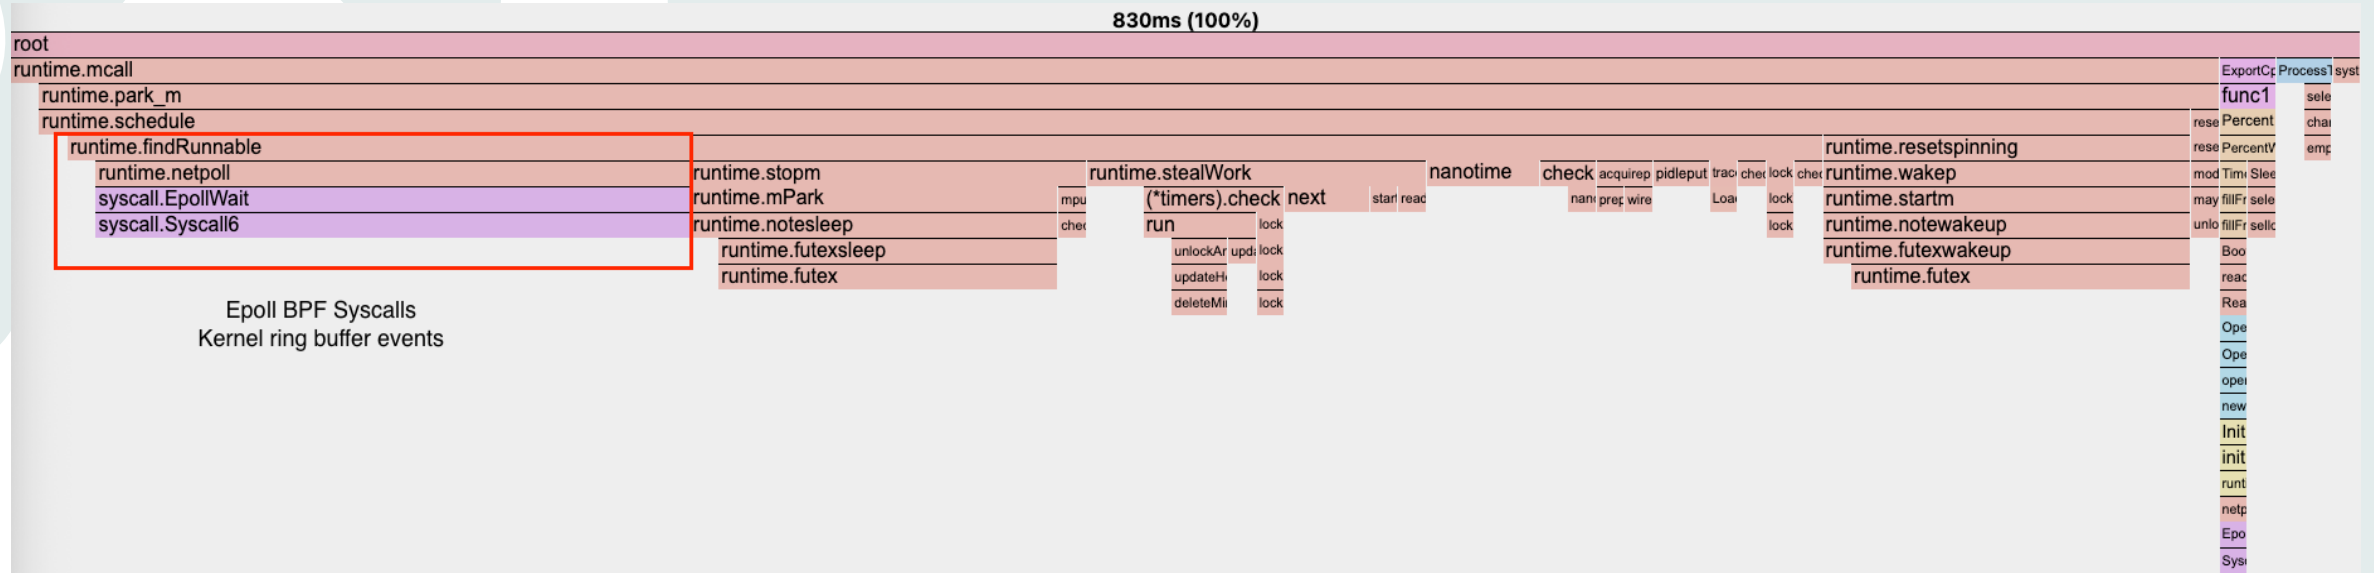
- Each Exfiltration Attempt




# Kernel eBPF Program DNS DPI Time (SKB Parsing)



# Resource Usage – eBPF Agent Flame Graph



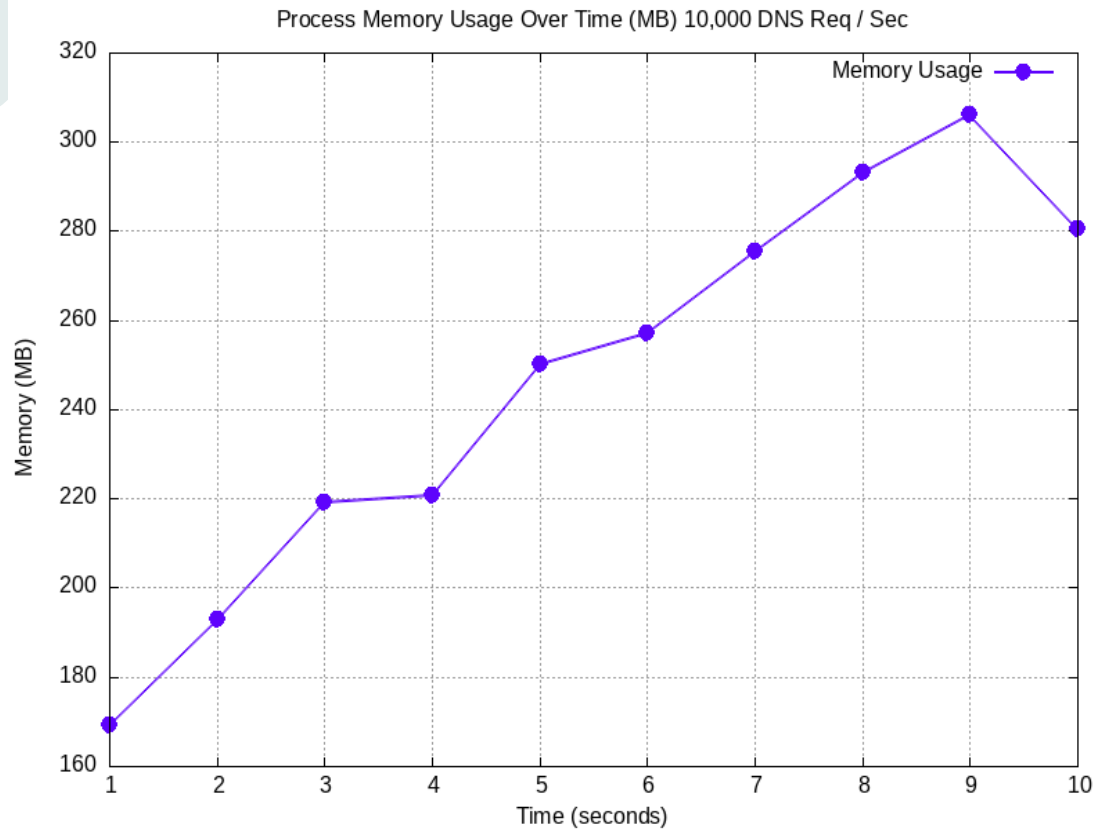
Idle Usage



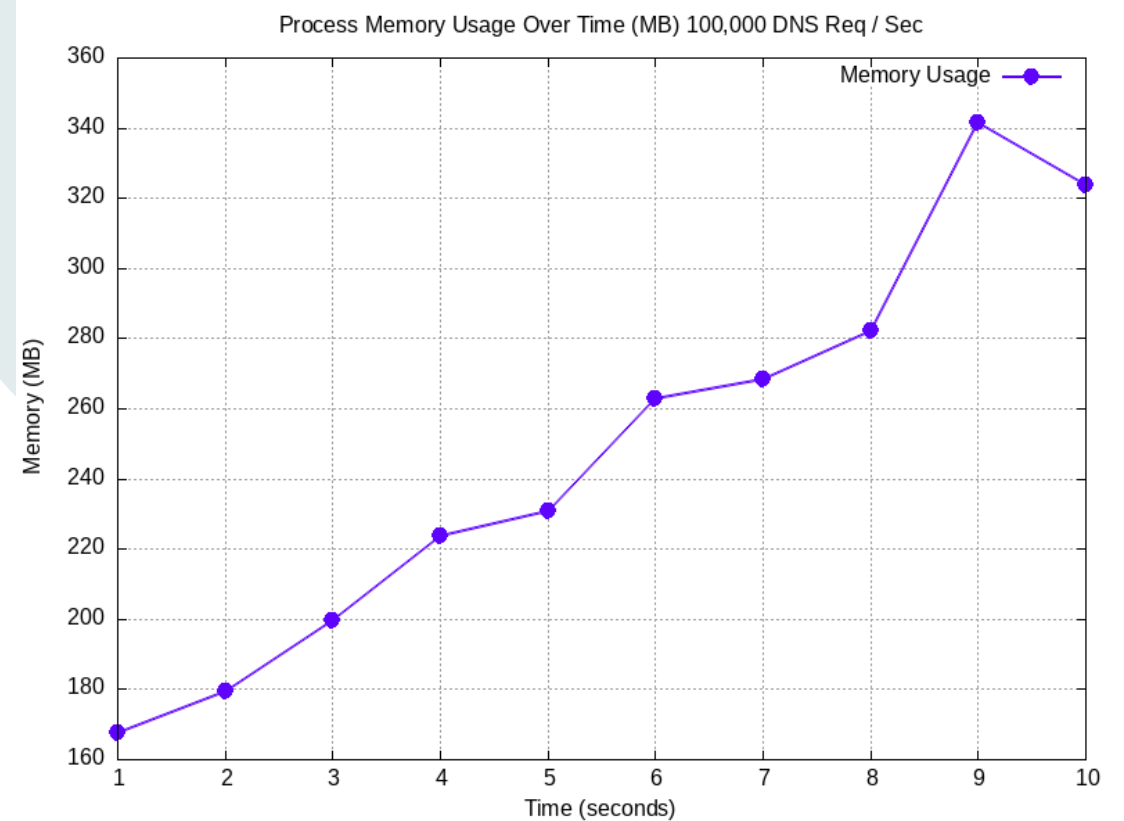
CPU Throttle Run



# Resource Usage - Memory

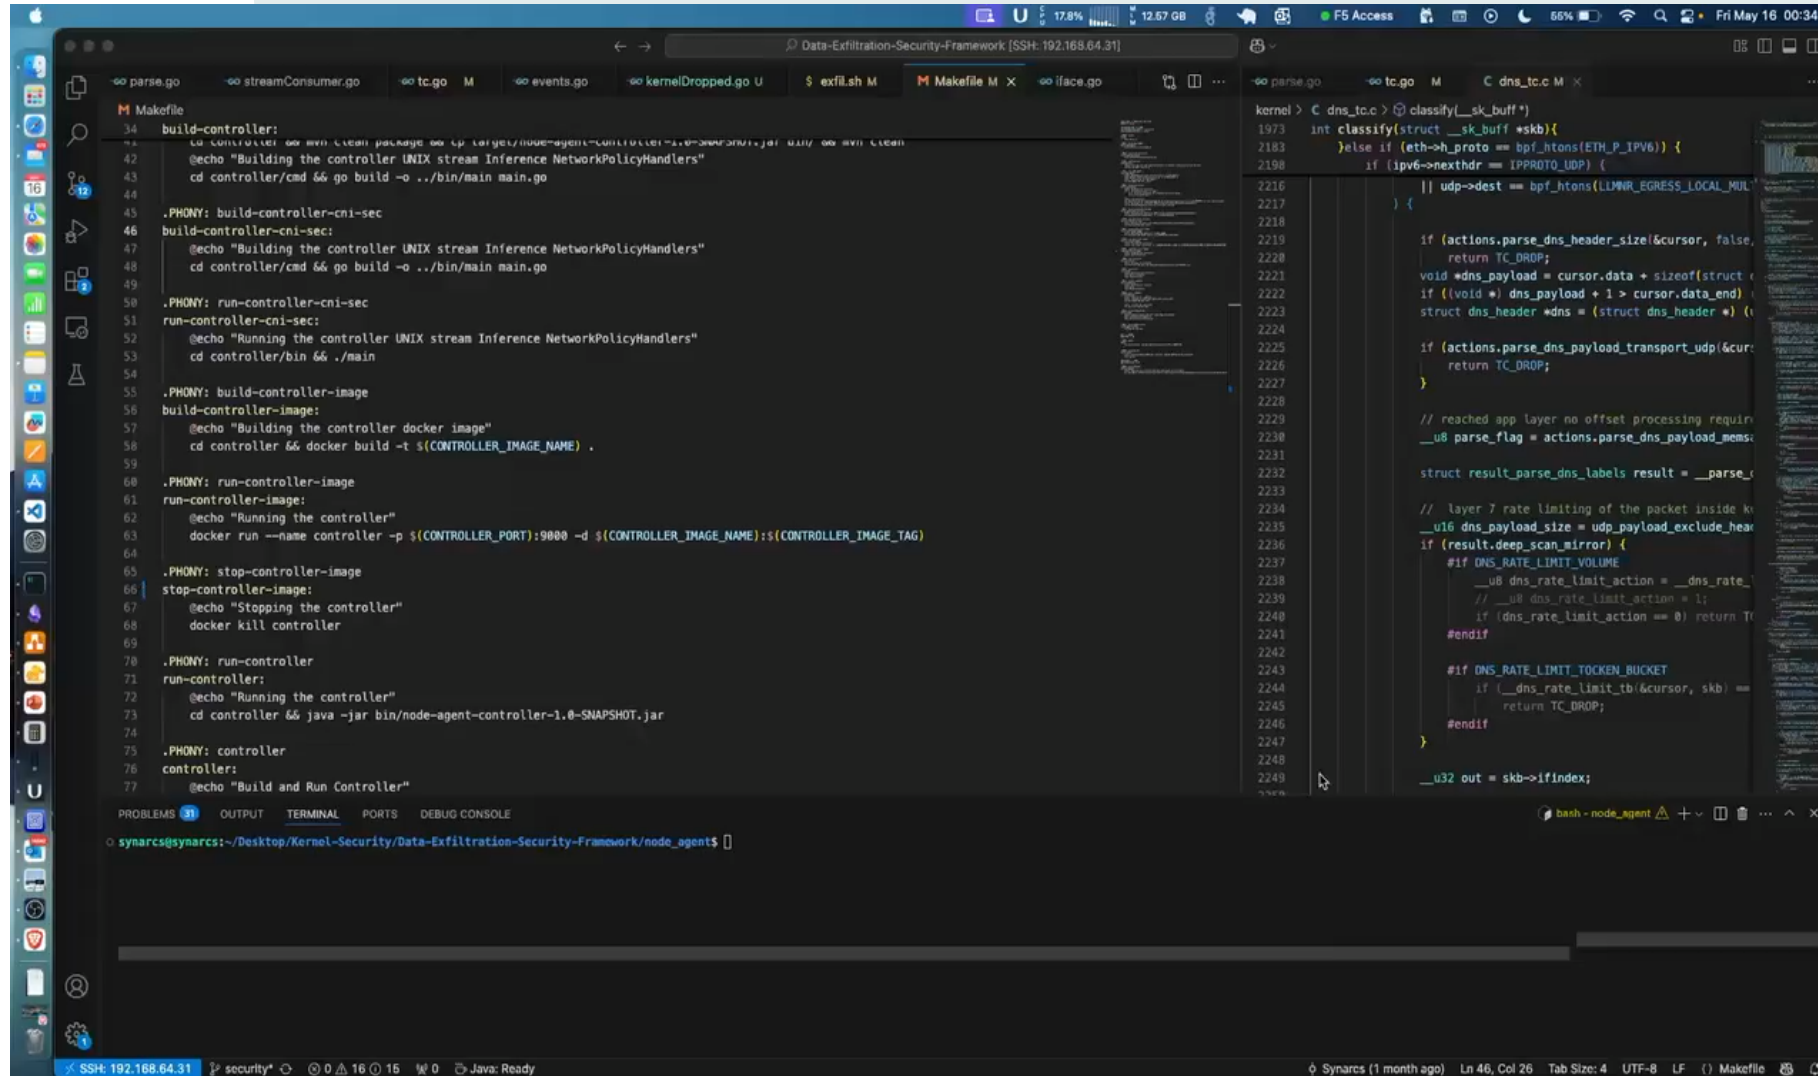


10,000 DNS Req / Sec



100,000 DNS Req / Sec

# Framework Security Strength : Demo



The screenshot displays a VS Code editor window with two files open: `Makefile` and `dns_tc.c`. The `Makefile` on the left contains build instructions for a controller, including targets for building the controller binary, running it, and building a Docker image. The `dns_tc.c` file on the right shows a C function `classify_skb` that processes network packets, including DNS headers and payloads, and applies rate limiting. The bottom status bar indicates the current file is `Makefile` at line 46, column 26.

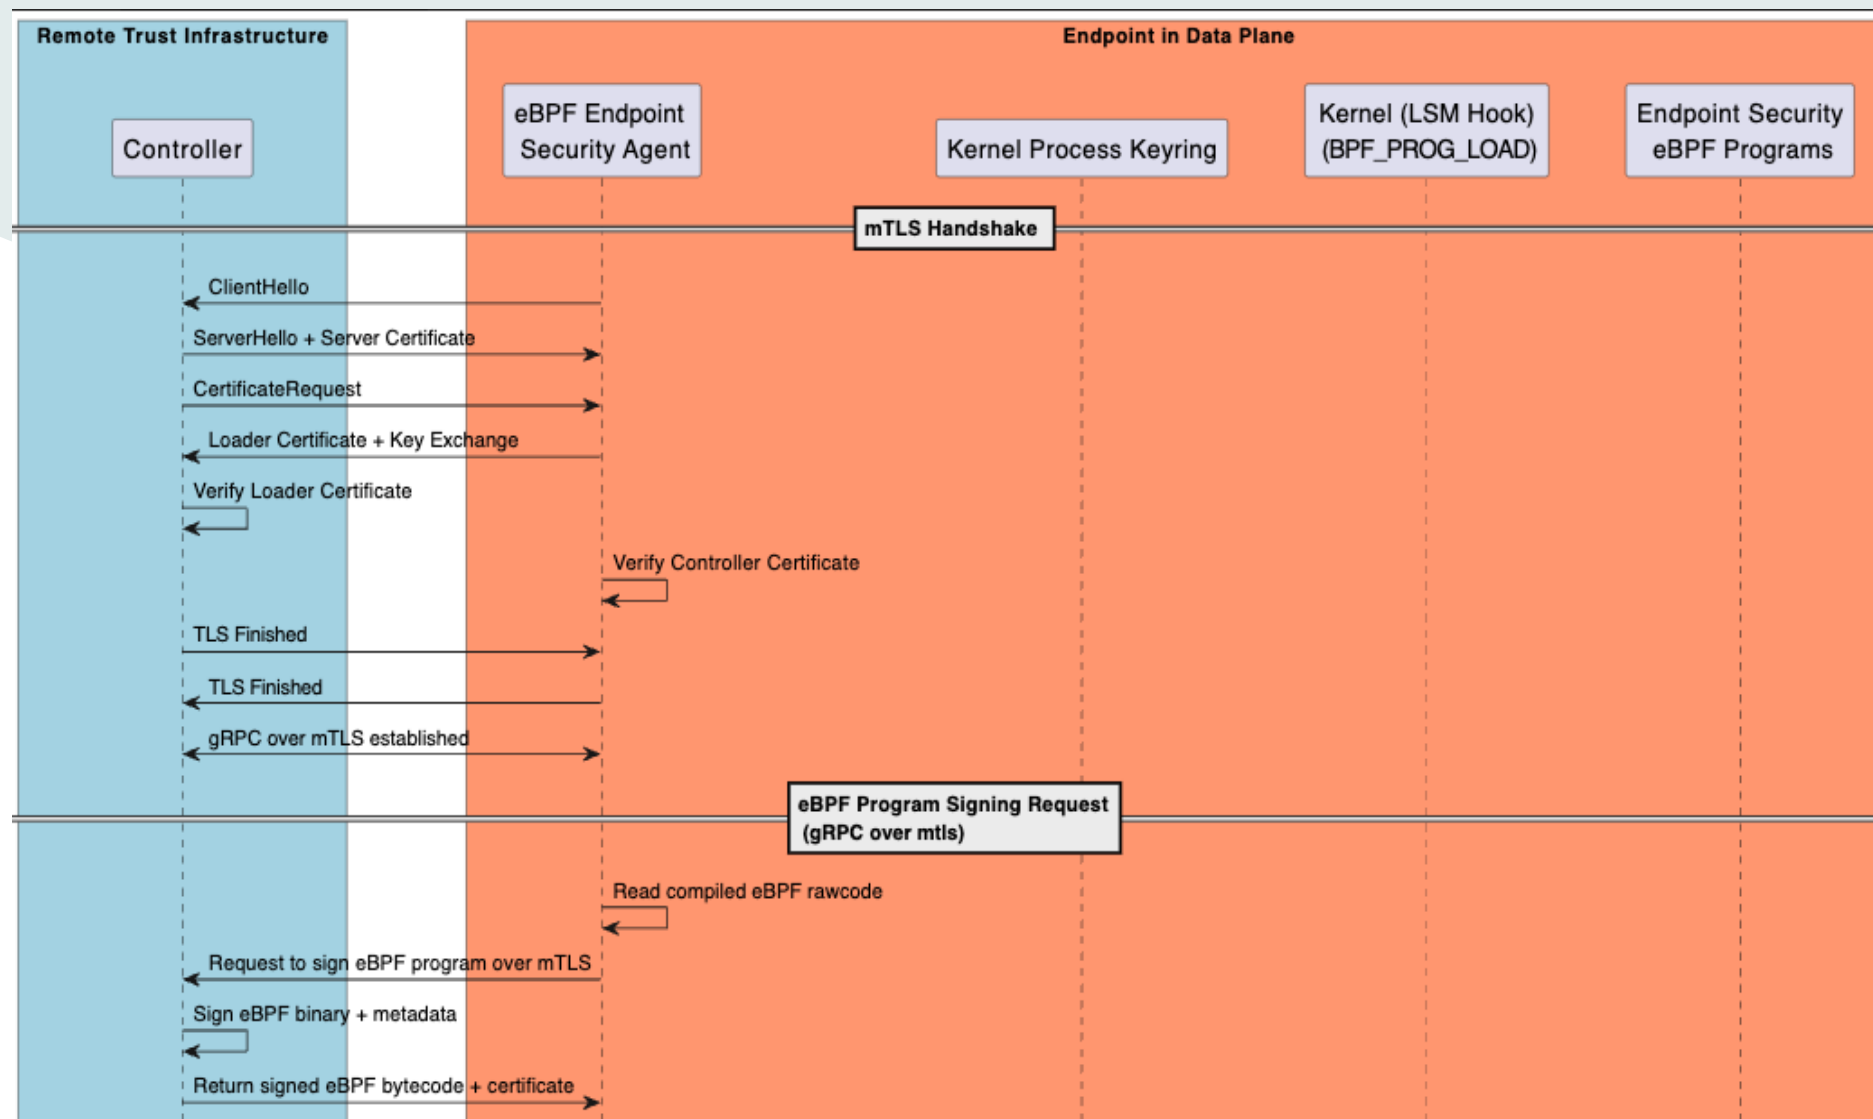
```
34 build-controller:
35     @echo "Building the controller UNIX stream Inference NetworkPolicyHandlers"
36     cd controller/cmd && go build -o ../bin/main main.go
37
38 .PHONY: build-controller-cni-sec
39 build-controller-cni-sec:
40     @echo "Building the controller UNIX stream Inference NetworkPolicyHandlers"
41     cd controller/cmd && go build -o ../bin/main main.go
42
43 .PHONY: run-controller-cni-sec
44 run-controller-cni-sec:
45     @echo "Running the controller UNIX stream Inference NetworkPolicyHandlers"
46     cd controller/bin && ./main
47
48 .PHONY: build-controller-image
49 build-controller-image:
50     @echo "Building the controller docker image"
51     cd controller && docker build -t $(CONTROLLER_IMAGE_NAME) .
52
53 .PHONY: run-controller-image
54 run-controller-image:
55     @echo "Running the controller"
56     docker run --name controller -p $(CONTROLLER_PORT):9000 -d $(CONTROLLER_IMAGE_NAME):$(CONTROLLER_IMAGE_TAG)
57
58 .PHONY: stop-controller-image
59 stop-controller-image:
60     @echo "Stopping the controller"
61     docker kill controller
62
63 .PHONY: run-controller
64 run-controller:
65     @echo "Running the controller"
66     cd controller && java -jar bin/node-agent-controller-1.0-SNAPSHOT.jar
67
68 .PHONY: controller
69 controller:
70     @echo "Build and Run Controller"
```

```
1973 int classify(struct __sk_buff *skb){
2183 }else if (eth->h_proto == bpf_htons(ETH_P_IPV6)) {
2198     if (ipv6->nexthdr == IPPROTO_UDP) {
2216         || udp->dest == bpf_htons(LLMNR_EGRESS_LOCAL_MUL"
2217     ) {
2218         if (actions.parse_dns_header_size&cursor, false,
2219             return TC_DROP;
2220         void *dns_payload = cursor.data + sizeof(struct {
2221             if ((void *) dns_payload + 1 > cursor.data_end)
2222             struct dns_header *dns = (struct dns_header *) (
2223             if (actions.parse_dns_payload_transport_udp(&cur:
2224                 return TC_DROP;
2225             }
2226             // reached app layer no offset processing requir
2227             __u8 parse_flag = actions.parse_dns_payload_mems:
2228             struct result_parse_dns_labels result = __parse(
2229             // layer 7 rate limiting of the packet inside k
2230             __u16 dns_payload_size = udp_payload_exclude_hear
2231             if (result.deep_scan_mirror) {
2232                 #if DNS_RATE_LIMIT_VOLUME
2233                 __u8 dns_rate_limit_action = __dns_rate_
2234                 // __u8 dns_rate_limit_action = 1;
2235                 if (dns_rate_limit_action == 0) return TC
2236                 #endif
2237                 #if DNS_RATE_LIMIT_TOKEN_BUCKET
2238                 if (__dns_rate_limit_tb(&cursor, skb) ==
2239                     return TC_DROP;
2240                 #endif
2241             }
2242             __u32 out = skb->ifindex;
```

# First Trust Chain (Endpoint-Agents – Cloud Trust Infrastructure)

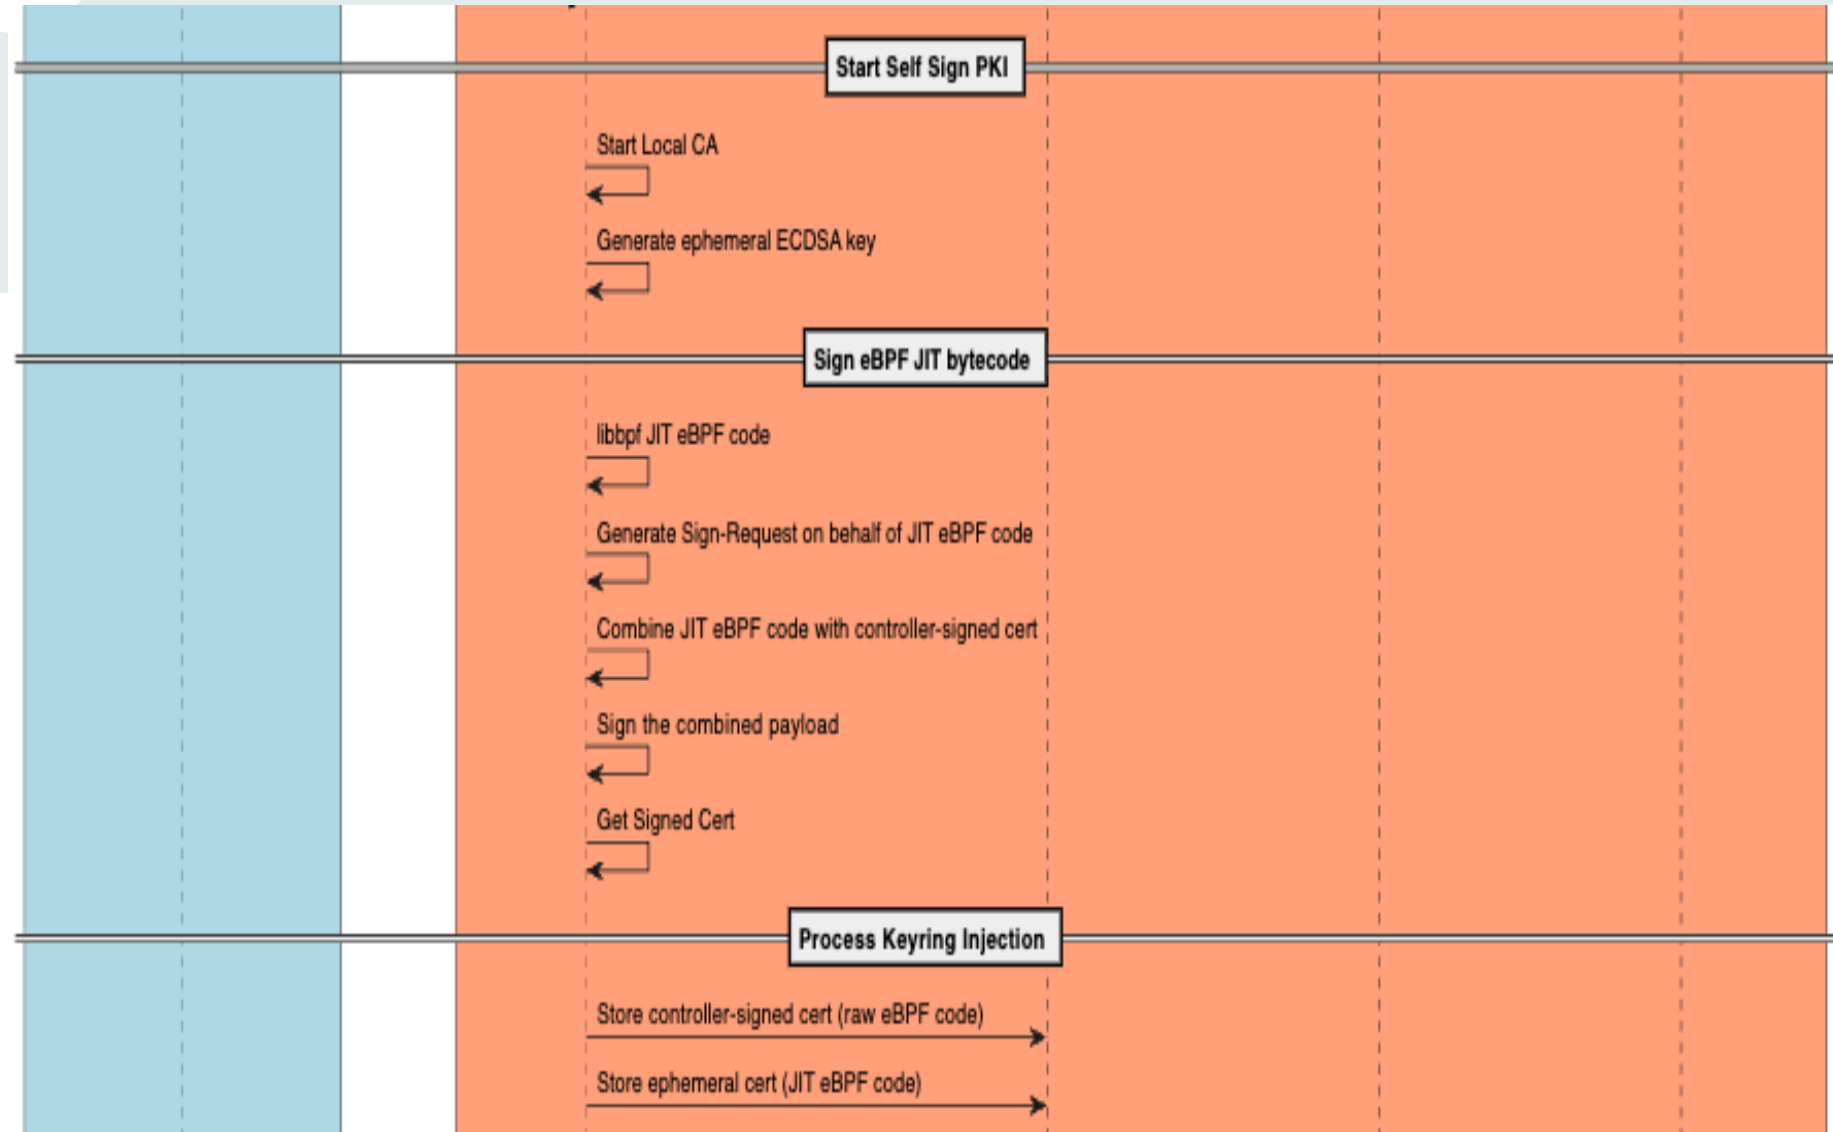
- **Endpoint Agents**
- **Controller (Remote PKI)**

- **mTLS-based identity handshake**
- **Remote-signed eBPF raw bytecode**



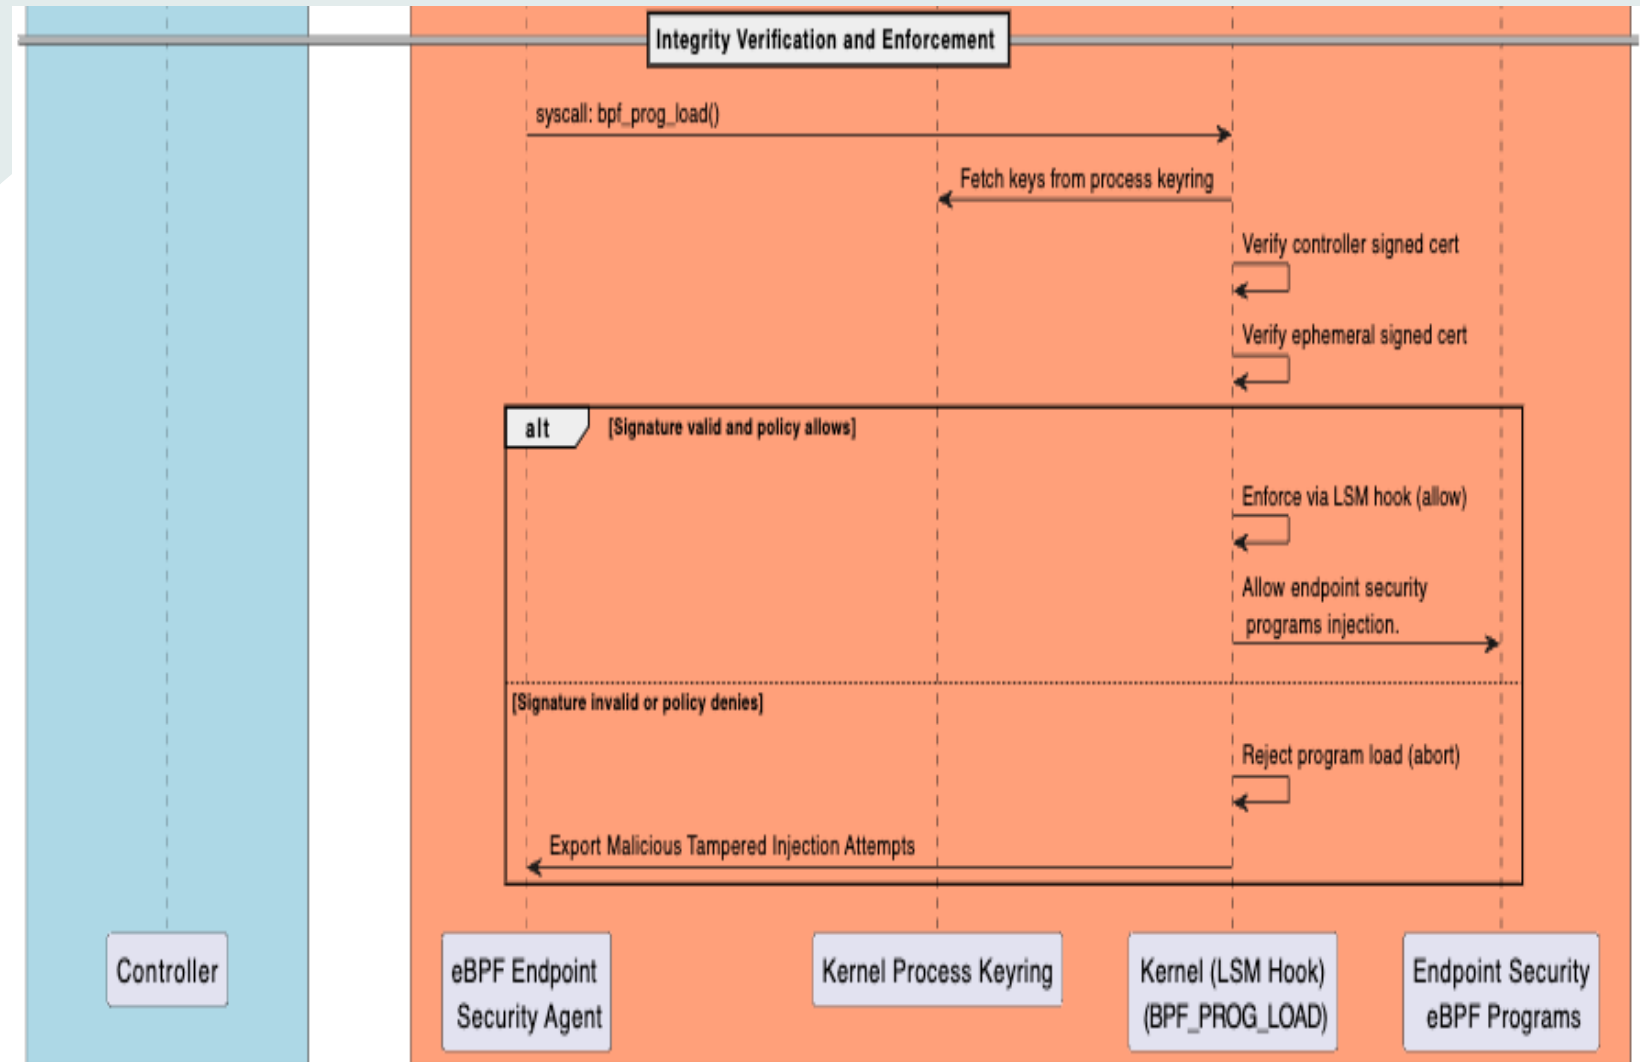
# Ephemeral Runtime Signing Bootstrap PKI (JIT bytecode integrity)

- **Ephemeral PKI bootstrapping**
- **JIT eBPF bytecode signing**
- **Controller-bound cert injection**
- **Keyring-based program identity**



# Second Trust Chain (Endpoint-Agents – Kernel Keyring)

- **Verify controller signed cert (raw eBPF bytecode)**
- **Verify ephemeral signed cert (JIT bytecode)**
- **Ensure 3-way strong integrity prevent eBPF program tampering.**



# Summary and Enhancements

- **Endless Security Enforcement Loop:** Kernel-enforced security strengthens EDR with real-time defense against advanced C2.
- **Dynamic in-kernel security enforcement through eBPF:** Reprogram kernel with proactive defense over advanced C2 mutating nature.
- **Extend Support for DNS-over-TCP:** Implement in-kernel eBPF-based detection for DNS-over-TCP replicating TCP state machine over kernel socket layer, paired with userspace DPI via Envoy proxy.
- **Add In-Kernel TLS Fingerprinting and Encrypted Tunnels:** Use eBPF for TLS fingerprinting (e.g., JA3/JA4) to detect DNS exfiltration over TLS (DOH), DNS over TLS, WireGuard.
- **Continuous Model Evolution:** Drift detection, online learning, and confidence-based live updates to maintain precision against emerging DNS obfuscation tactics.
- **Cloud Native Security:**
  - Dynamic L3/L7 security enforcement over cloud Vnet's / VPC via dynamic blacklist's NACL's.





# Discussions and Q&A

## Code:

<https://github.com/Synarcs/DNSObelisk>

## WhitePaper:

[https://github.com/Synarcs/DNSObelisk\\_Report](https://github.com/Synarcs/DNSObelisk_Report)

### Endpoint Agent LOC

	Language	LOC
Kernel (eBPF)	C	3314
Userspace	Go, Python	8862

### Rest Framework LOC

	Language	LOC
Userspace	Go, Java, Lua	1991