

# DNS Data Exfiltration Prevention: Kernel-Enforced Endpoint Security

*Scalable Framework to Disrupt  
DNS C2 and Tunneling*

Vedang Parasnis

University of Washington Bothell








# What is Data Exfiltration


[[Singamaneni et al.](#)]

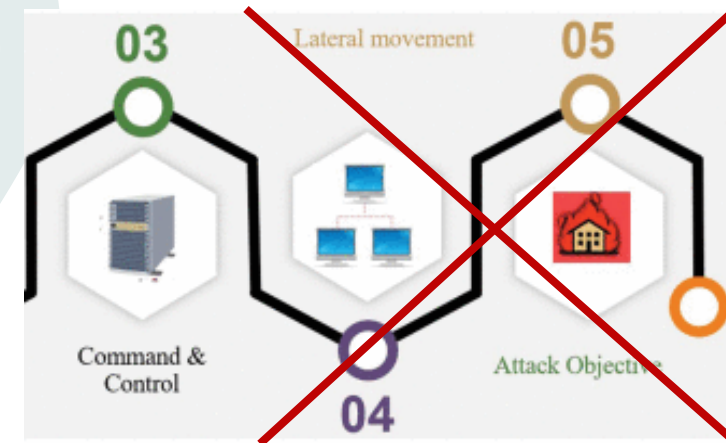
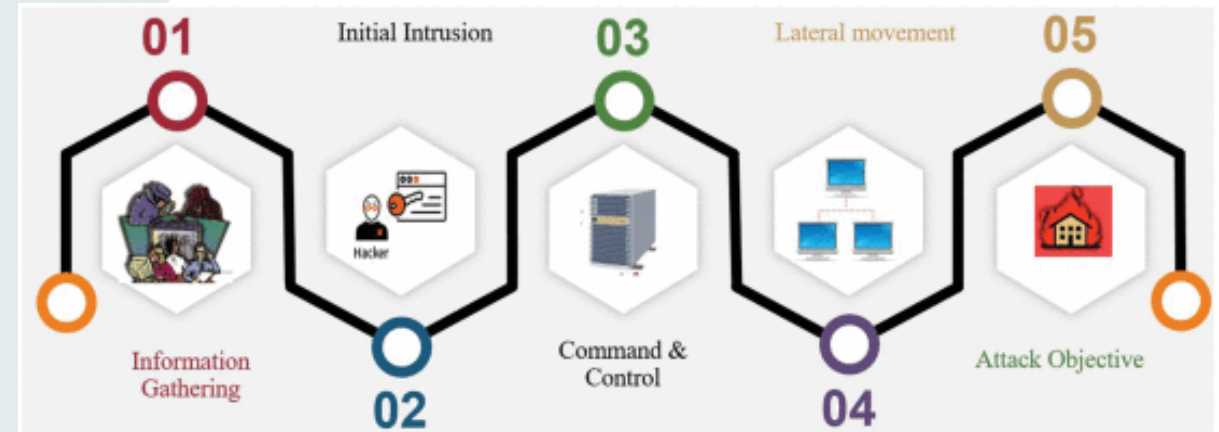
**Definition:** Unauthorized extraction or transmission of sensitive data from a system

## Attack Lifecycle

-  Information Reconnaissance
-  Initial Intrusion / Infiltration
-  **Command and Control**
-  Lateral Movement
-  Command Execution and Data Breaches

## Core Defense Strategy

-  Endpoint Security (EDR / XDR)

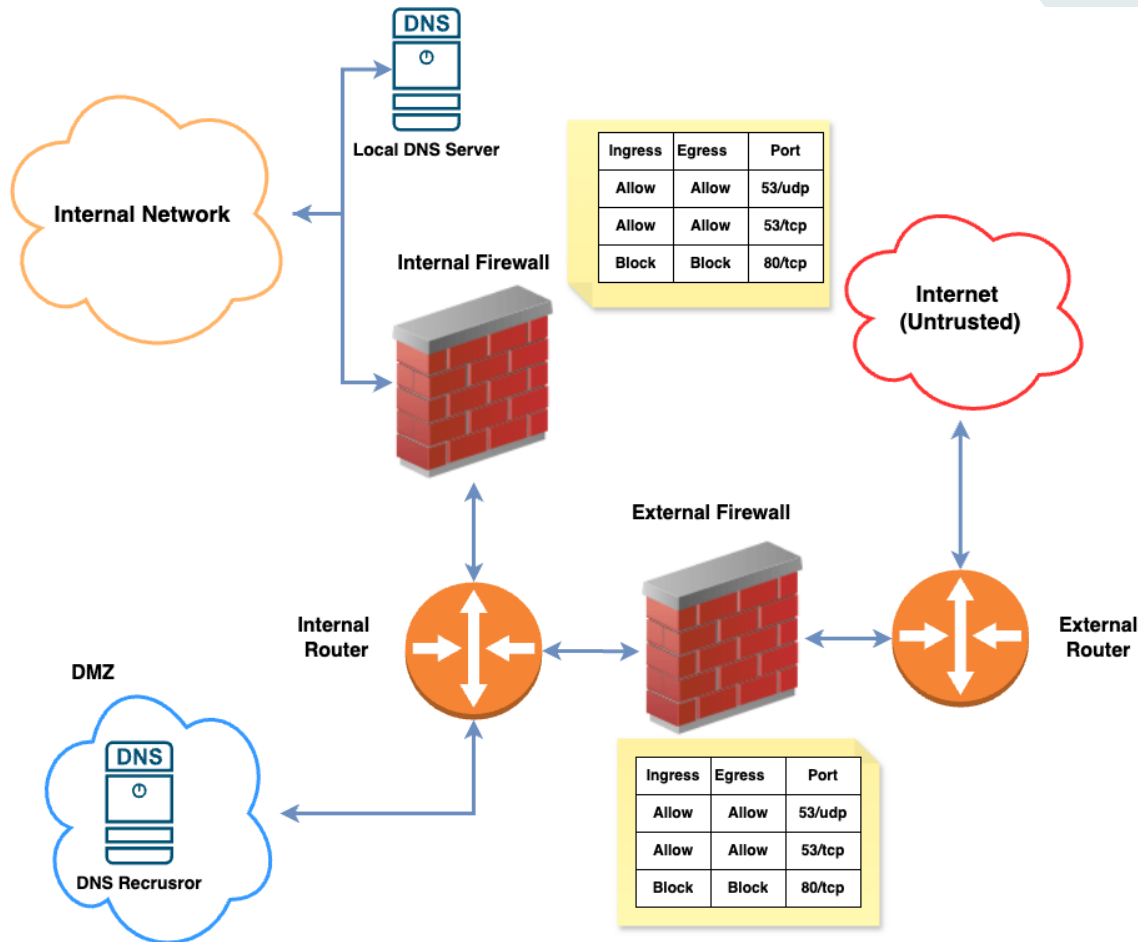


# Why DNS is a Blind Spot

**Unencrypted by default** – Allows attackers to hide malicious payloads in plain sight.

**Rarely monitored deeply** – DNS logs are often ignored, giving adversaries a free channel.

**Firewall blindspot** – DNS ports (53 UDP/TCP) stay open, bypassing most traditional defenses.



# DNS Exfiltration Attack Methods



**DNS C2:** Enables stealthy remote control of compromised systems by tunneling commands through DNS traffic

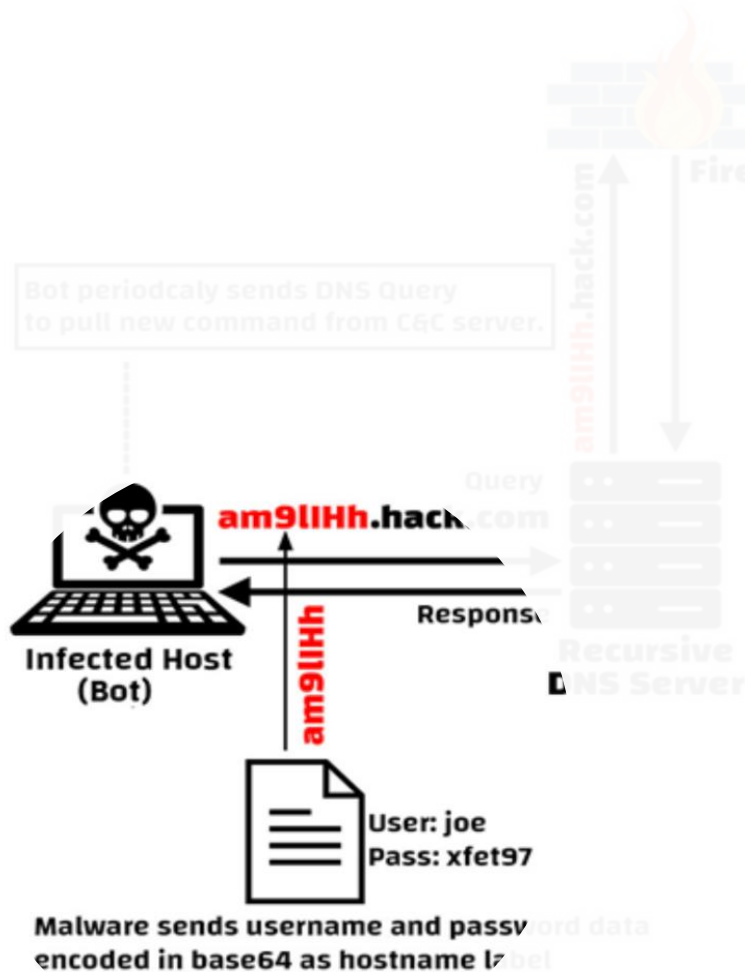


**DNS Tunneling:** Abuses DNS protocol to bypass network controls and exfiltrate data or maintain covert communication channels



**DNS Raw Exfiltration:** Leaks sensitive data directly within raw DNS queries, evading traditional detection mechanisms

# DNS Data Exfiltration



## Remote Code Execution (RCE)

- Shell code exploits
- Script executions, File corruptions
- Process Side channeling exploits
- Example: **Sliver C2, Hexane, APT29 (Cozy Bear).**

## Persistent Backdoors

- Deployment rootkits, ransomwares
- Example: **Turla group**

## Network Pivoting (Port Forwarding)

- Compromised machines act as proxies to reach deeper into private infrastructure
- Example: **Cobalt Strike, Hexane, DNSSystem**

# Existing Approaches

- **Active Analysis**
  - **DNS Exfiltration Security as Middleware**
    - Palo Alto Precision Guard AI Security
    - Infoblox DNS exfiltration security
- **Passive Analysis**
  - Anomaly Detection [[Bilge et al.](#)]
  - Threat Signatures, Domain Reputation scoring [[Antonakakis et al.](#)]

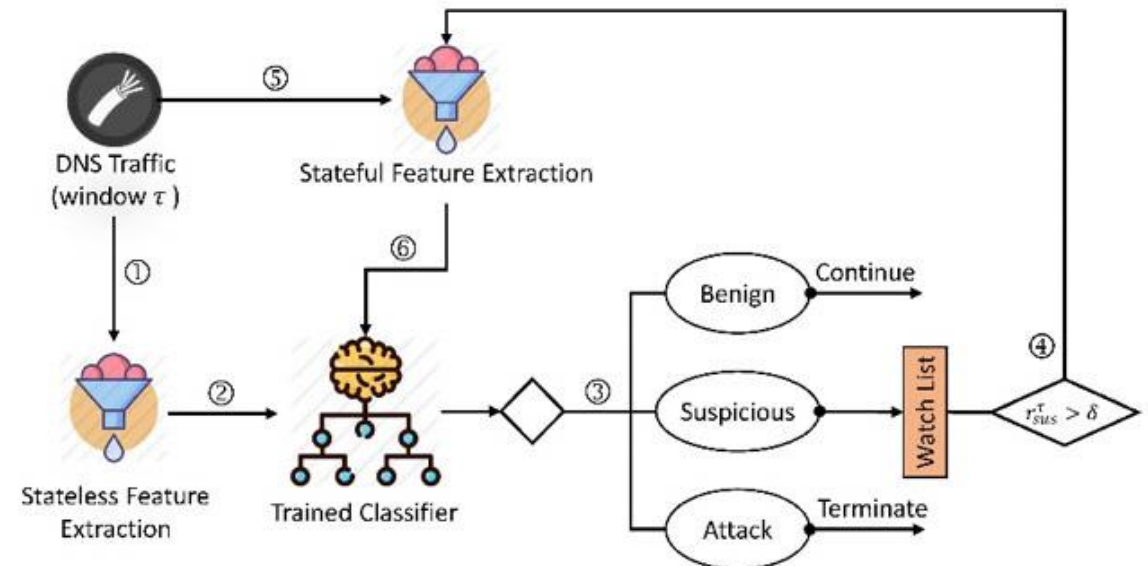


# Existing Approaches – Passive Analysis

- **Anomaly Detection:**
  - **Traffic Behavior Analysis**
    - DNS Passive Traffic Volume Analysis
    - DNS Passive Traffic timing Statistical Analysis
  - **Machine Learning-based Threat Intelligence**
    - Uses machine learning models to identify attack behavior.
- **Threat Signatures:**
  - **DNS Domain Scoring**
  - **Malicious domain signature**

Stateless Features – Lexical Analysis

Stateful Features – Statistical Analysis



[[Samaneh et al.](#), [Jawad et al.](#)]

# Issues with current approaches

- **Slow Detection → High Dwell Time → More Damage**
- **Extremely slow to Advanced C2 Attacks**
- **Kernel Encapsulated Traffic**
- **Dynamic Threat Patterns:**
  - Varying Throughput
  - Slow and Stealthy Rate
- **Centralized monitoring and analysis systems don't scale**
- **Ineffective over IP Masquerading & Domain Generation Algorithms**

## **Solution:**

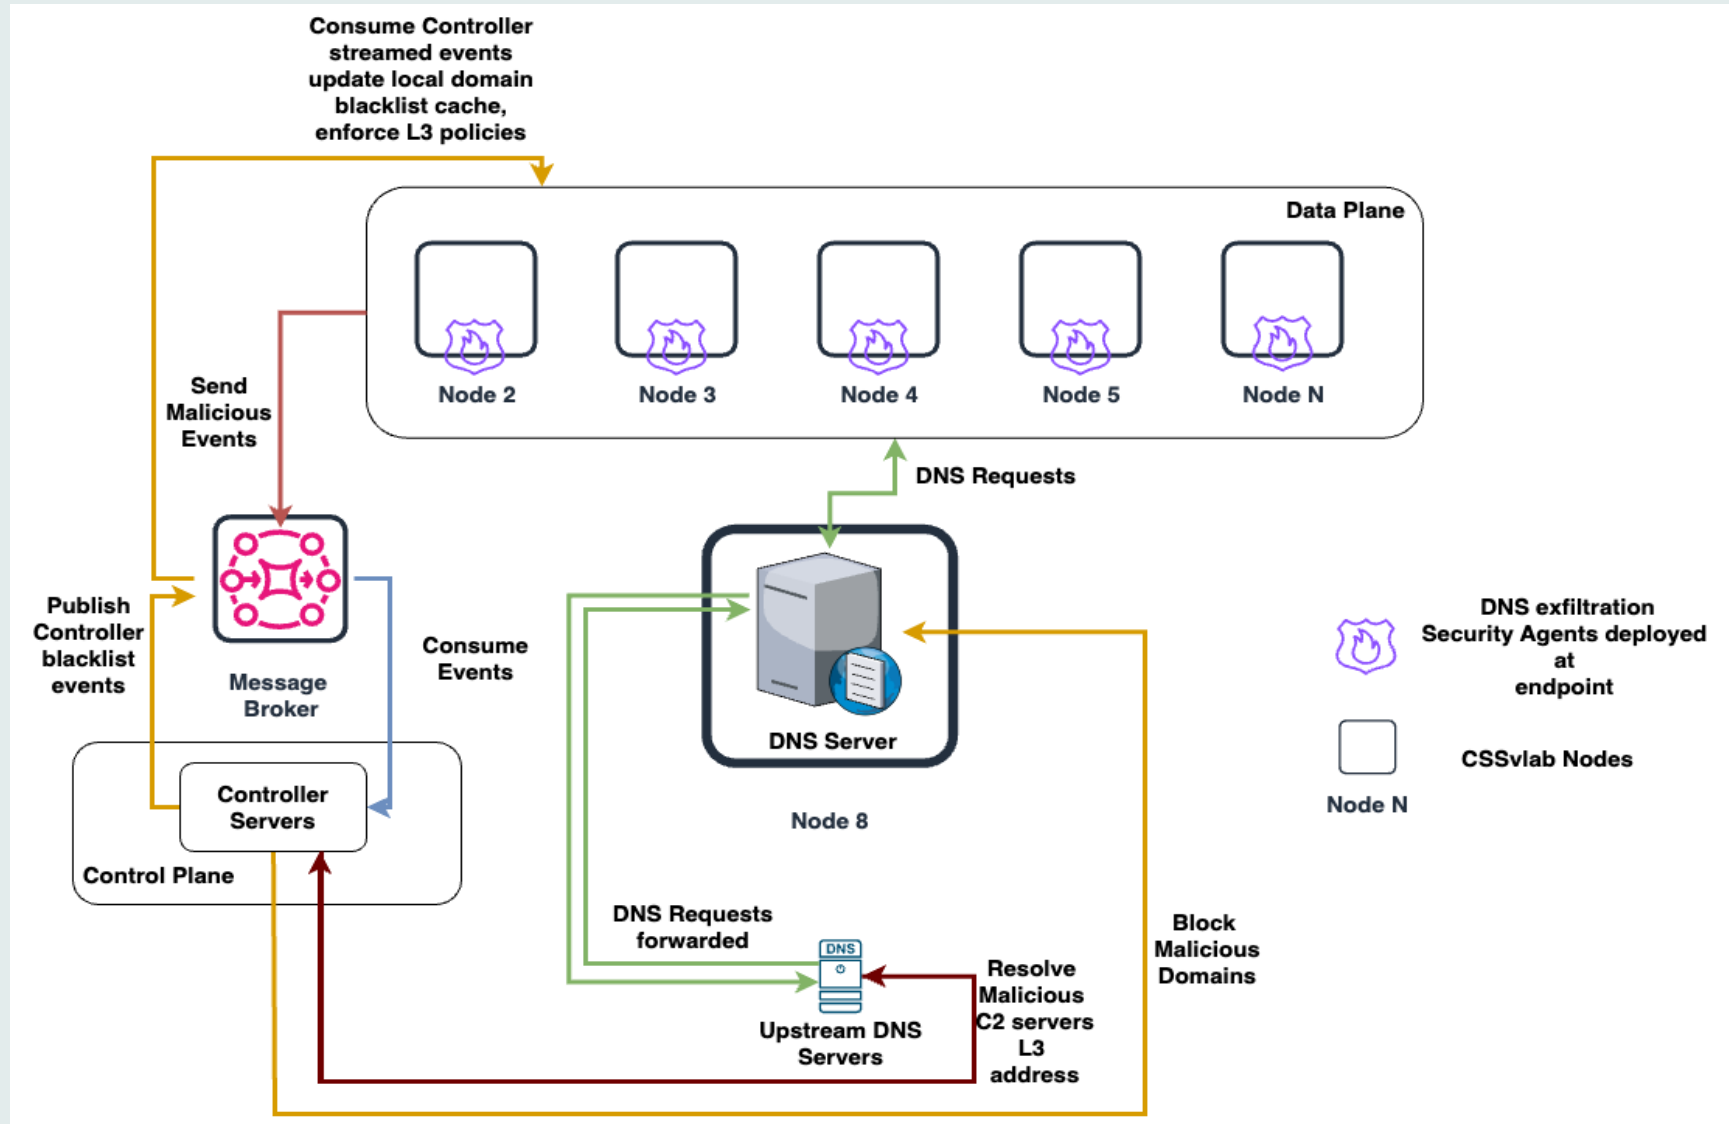
**Real-time, proactive enforcement at Ring 0 — inside the kernel, where no userland evasion can hide.**



# Security Framework Architecture

## Architecture Components

- **Data Plane**
  - eBPF endpoint agents
- **Control Plane**
  - Controller Servers
- **Infrastructure**
  - DNS Servers
  - Apache Kafka



# Security Framework Goals

## **Real-Time DNS Exfiltration Prevention**

Implement in-kernel deep packet inspection and enforcement to block all forms of DNS exfiltration channels

## **AI-Assisted Threat Detection**

Use deep learning in userspace to detect advanced obfuscated exfiltration payloads with high accuracy aiding kernel enforcements.

## **Dynamic Cross-Layer Policy Enforcement**

Enforce in-kernel L3 network policies adaptively and domain blacklisting on DNS server to combat DGA.

## **Malicious Process Aware Active Defense**

Instantly detect and kill implants, preventing lateral movement and further damage.

## **Scalable Multi-Cloud Deployment**

Ensure framework's horizontal scales for real-world production cloud environments..


# Broader Impact and Applicability

- **Cloud Providers & HyperScalers**
  - Strengthens DNS-layer security in managed services.
  - Examples: AWS Route 53, Google Cloud DNS, Azure DNS.
- **National Security & Defense**
  - Disrupts advanced malware APT groups alive using DNS-based C2 channels.
  - Examples: Turla Venom, Skitnet, Lazarus, OilRig, Hexane.
- **Regulated Enterprises (Finance, Healthcare)**
  - Augments DLP capabilities over DNS for private cloud and on-premise environments.
  - Examples: Financial institutions, healthcare networks.
- **Security Vendors (EDR/XDR/DNS Security)**
  - Integrates as a modular addon to extend EDR/XDR threat prevention at the DNS level.
  - Examples: CrowdStrike Falcon, Cisco HyperShield, Palo Alto Precision AI, Broadcom Carbon Black.

# eBPF – Extended Berkley Packet Filter

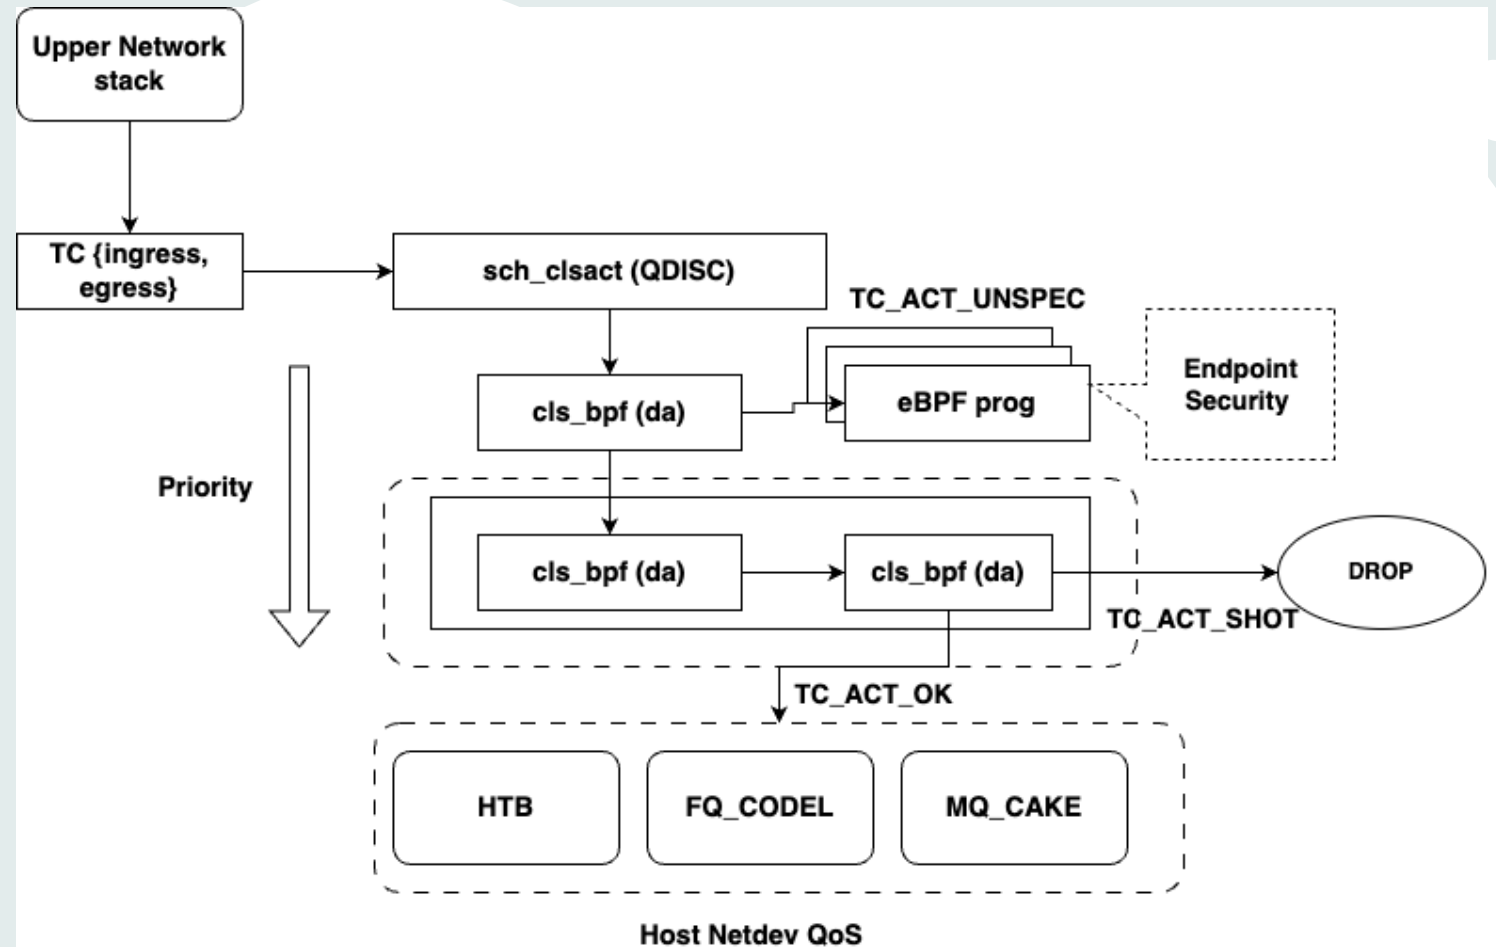
- Reprogram the Linux kernel in safe way:
- Modern way to write kernel modules

1. Runs BPF virtual machine inside kernel
2. Custom BPF bytecode
3. Uses 512 bytes of stack
4. eBPF Maps as heap
5. CPU architecture agnostic



# Linux Kernel Network Stack

- Sockets
- TCP/IP Stack
- Netfilter
- **Traffic Control (QoS)**
- Network Drivers



[*Jamal Salim et al.*, *Daniel Borkmann et al.*]

# Kernel Enforced Endpoint Security

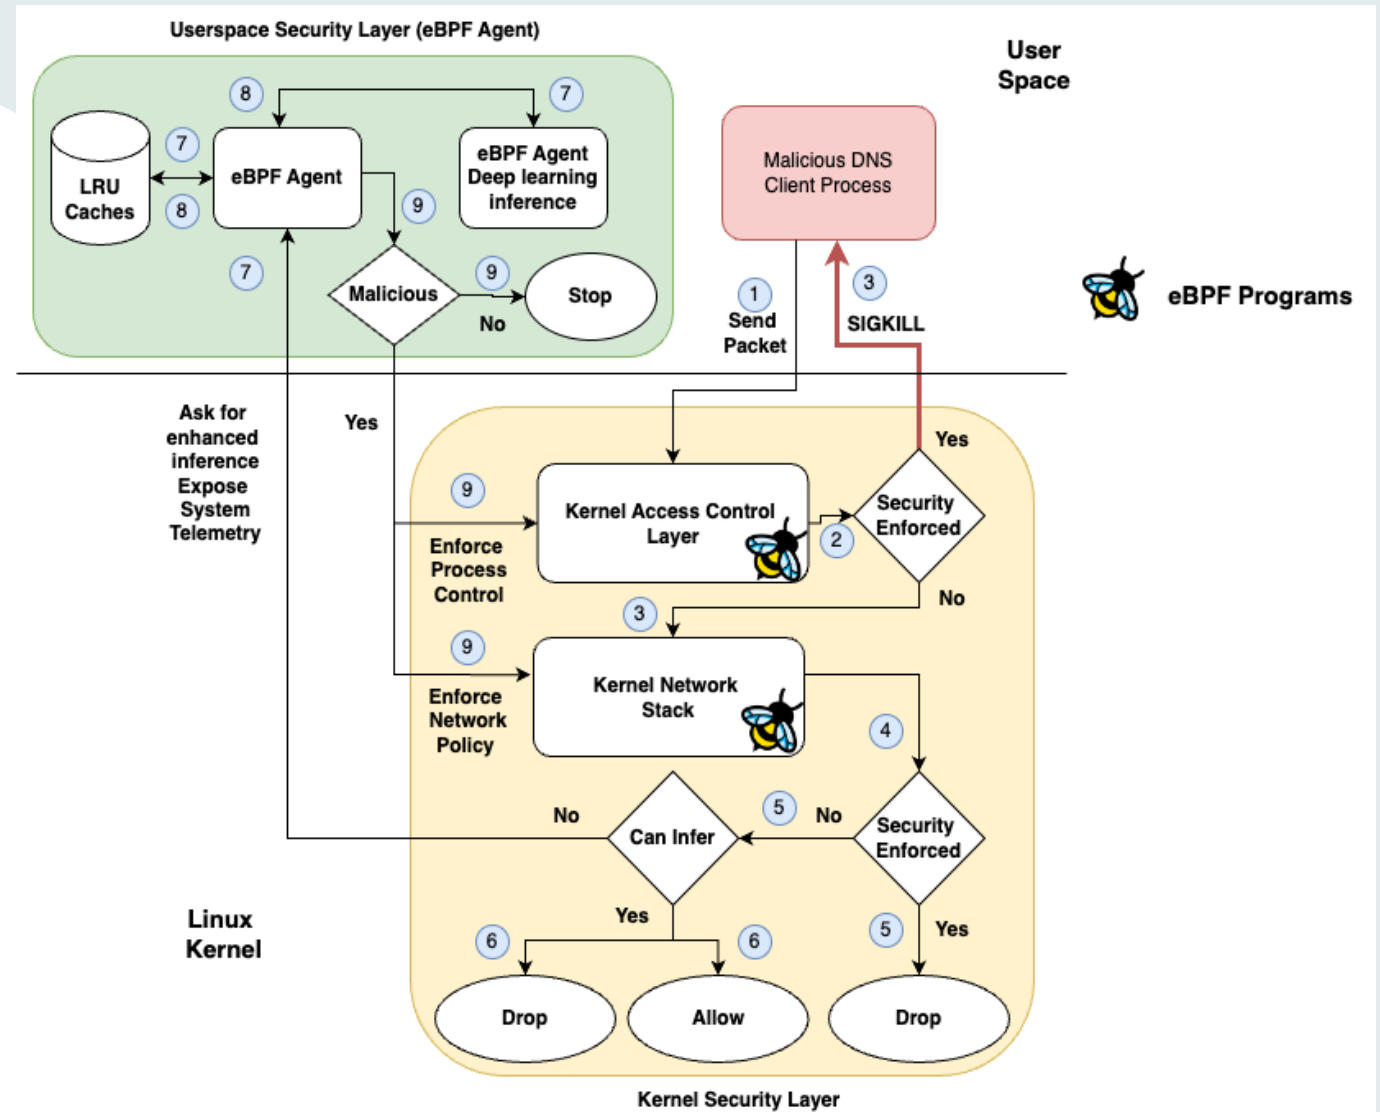
## Agent based Endpoint Security

### Userspace

- eBPF Userspace Bindings
- ONNX Quantized Deep Learning Model

### Linux Kernel

- Inference Unix Domain Sockets
- Kernel Network Stack (eBPF)
  - Socket Layer
  - Traffic Control
- Kernel Security Layer (eBPF)
  - Kernel Security Modules
  - Kernel Syscall






# eBPF Agent Operations Modes

**eBPF Agents in Data Plane handle DNS exfiltration over UDP**

Mode	Goal	Requirement	Security Enforcement Process
Strict Enforcement Active Mode	Kill C2 Implants, ensure zero data loss and C2 command execution	DNS Traffic over UDP ports (53, 5353, 5355), for encapsulated and non-encapsulated traffic.	<ul style="list-style-type: none"><li>• Live Redirects suspicious DNS packets to Userspace</li><li>• Trace malicious process exfiltration count and terminates it.</li></ul>
Process-Aware Adaptive Passive Mode	Kill C2 Implants, ensure negligible data loss and minimal C2 command execution	DNS Traffic over random UDP ports.	<ul style="list-style-type: none"><li>• Allow suspicious traffic passthrough</li><li>• In Kernel start threat hunting process tied to malicious DNS packets</li><li>• Trace malicious process and terminates it</li></ul>

# ONNX DNN Model

- Model Architecture ONNX Graph
- Sample Malicious Data
- Features



Feature	Description
---------	-------------

Malicious Exfiltrated data DNS queries	
381c018e3f5d05b78e3f6a026381e0f3476c066e8017be6ba9f5a9d758ef.d04bc3e0fc58e5a2401da590f3ee268a6af637eaafd210e58060a41082dc.92d594840bcb32a6500f39248db646e4e602f8547294692d83a4b4680223.b4d0ce0ec94abc9b6821cea90561aac558a6ba30b53e6b.bleed.io	
ae8c018e3f235392a20ca002649bd124bb6b506ba0771986720cbb1ad2e2.d59ca990aaa3eb1c580f5fb16d3b59d7eeb142458c8c54199c56e87b751c.69bbf57db184d263ed85a5ba5c9281ba327646f5638587016c9e0aa7b9b8.af182352de5de5b76a32242f04428b7d01b9a6d7999eb3.bleed.io7el4BGh376549344247687c217c3030393739363038373833303765353.bleed.io7el4BGh6a70677c217c52454749535445527c217c61343266363038366.bleed.io	
sebubx76xk4erpp3rwehoo3ubmbqeaqbaeq.a.e.e5.sk	
4az3kiecotwu3okbtvfm7pdpcabqeaqbaeq.a.e.e5.sk	
<b>B (1)</b> unk__6x1	<b>B (1)</b> unk__6x1
<b>B (1)</b> unk__6x1	<b>B (1)</b> unk__6x1

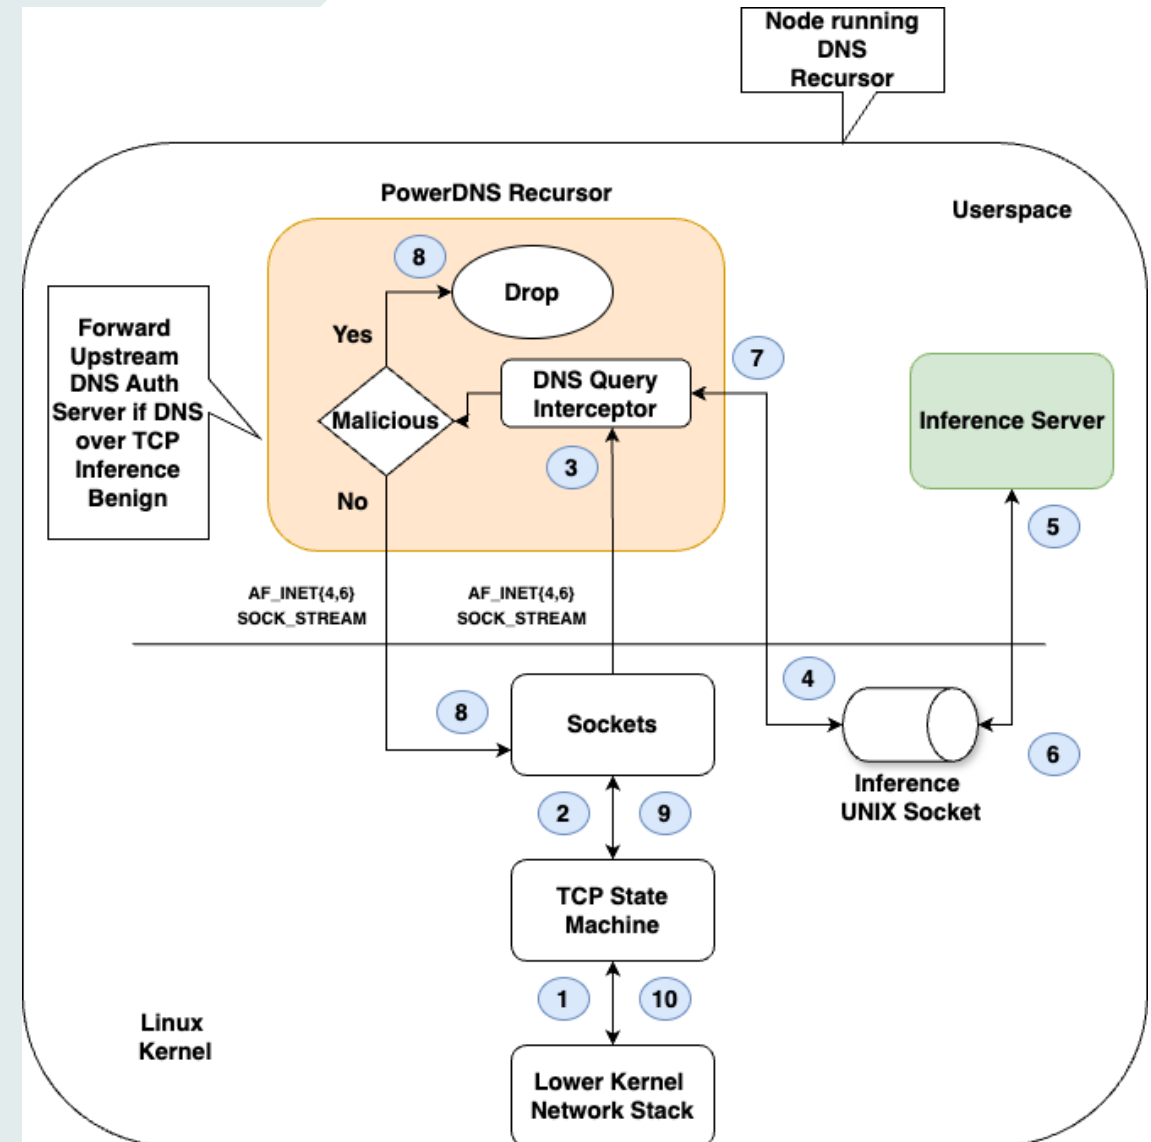
# Datasets

Dataset Type	Source / Characteristics	Size	Primary Goal
Trusted Benign Cache	Top 1M Cisco Second-Level Domains (SLDs)	1 Million	Reduce inference on known-good traffic.
ISP-Captured DNS	Live-sniffed ISP DNS traffic [ <a href="#">Ziza et al.</a> ]	50 Million	Provide real-world benign & malicious baseline.
Synthetic Exfiltration	Custom-generated (DET, DNSCat2, Sliver, Nuages, Custom Scripts, etc.);	2.4 Million	Malicious samples use varied obfuscation across file formats
Final Combined Dataset	Synthetically formed	3.8 millions	Balanced dataset w/ obfuscated payloads across file formats

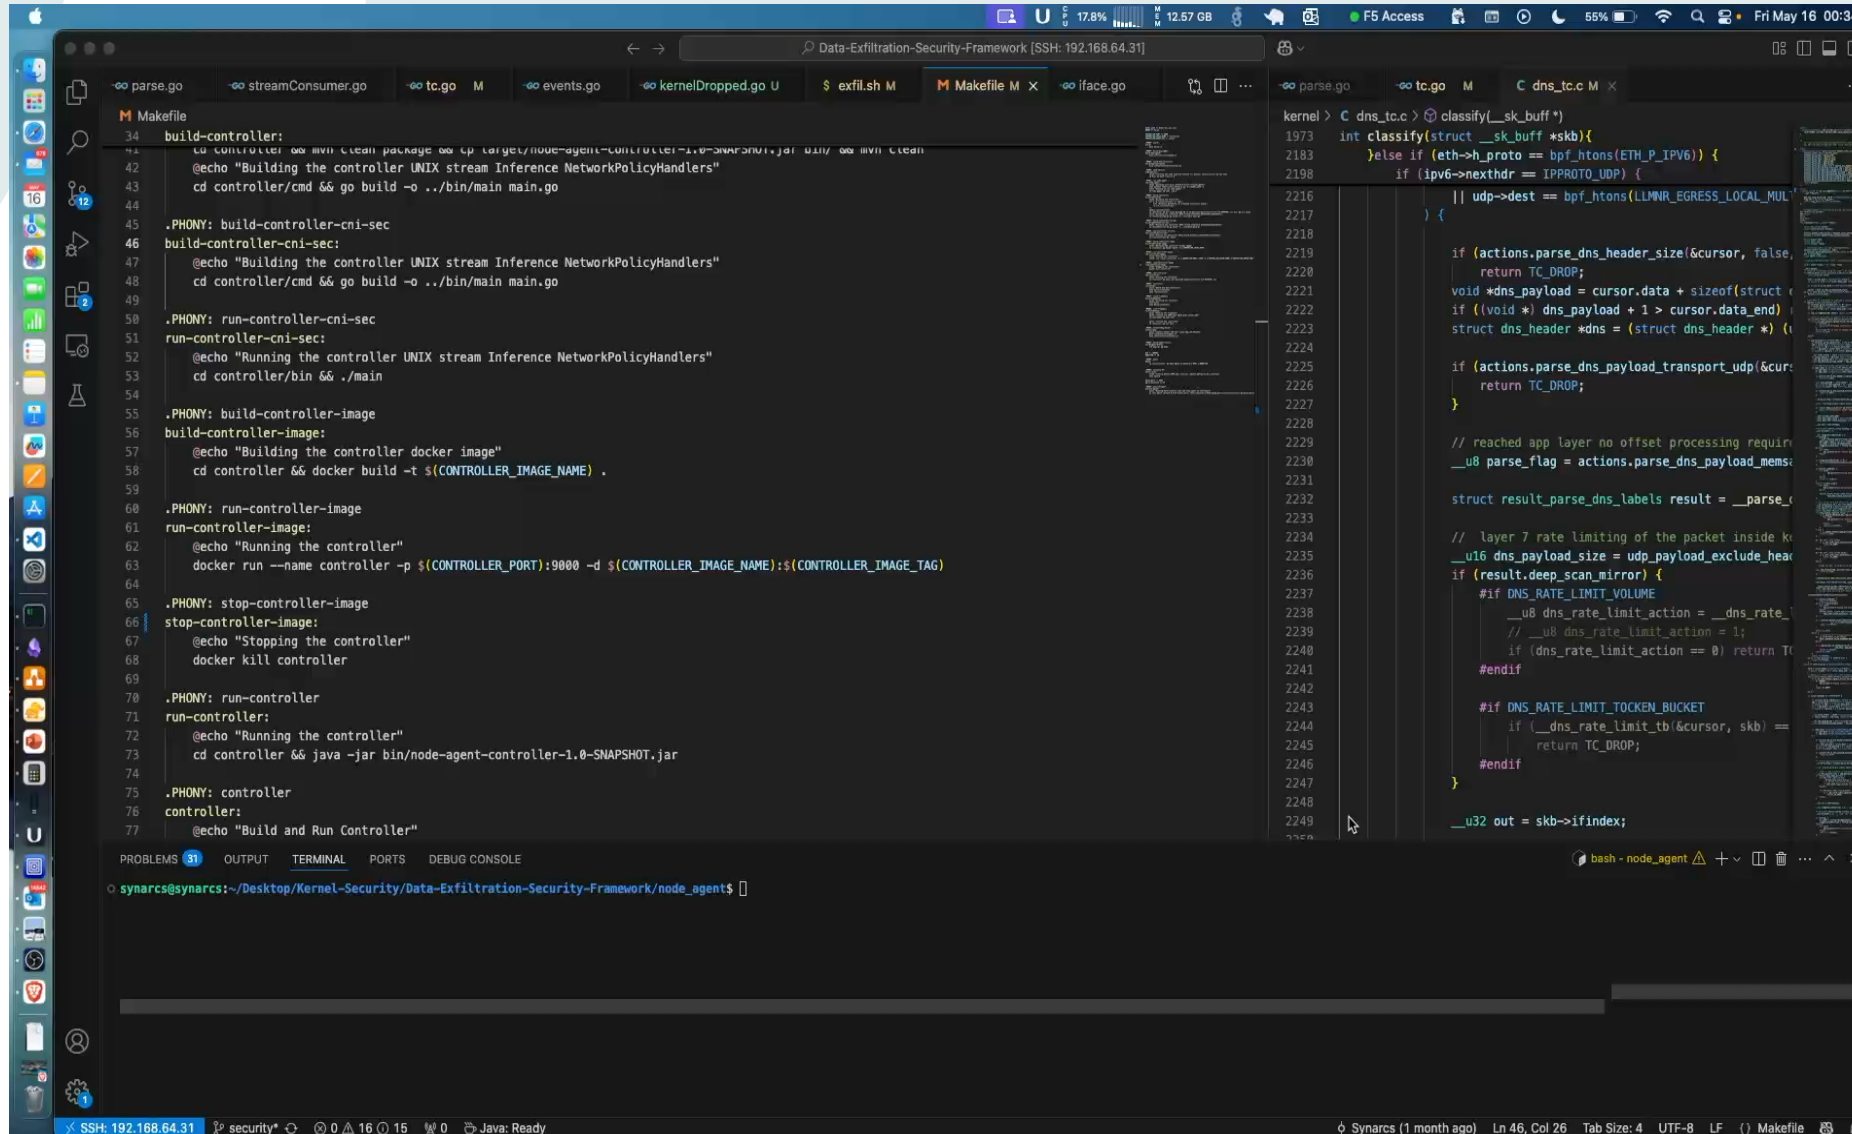
# Design Adjustment

## Prevent DNS Exfiltration over TCP

- Runs on
  - PowerDNS Recursor
- Relies on
  - PowerDNS recursor Query Interceptors
  - Inference UNIX domain sockets



# Framework Security Strength



The screenshot displays a terminal window with a dark theme. The top bar shows system information: 17.8% battery, 12.57 GB memory, and the date Fri May 16 00:34. The terminal is connected to a remote host via SSH (192.168.64.31). The main window is divided into two panes. The left pane shows a Makefile with targets for building and running a controller. The right pane shows C code for a kernel module, specifically a function to classify network packets. The bottom status bar indicates the current directory is ~/Desktop/Kernel-Security/Data-Exfiltration-Security-Framework/node\_agents and the terminal is running a bash shell.

```
34 build-controller:
35     @echo "Building the controller UNIX stream Inference NetworkPolicyHandlers"
36     cd controller/cmd && go build -o ../bin/main main.go
37
38 .PHONY: build-controller-cni-sec
39 build-controller-cni-sec:
40     @echo "Building the controller UNIX stream Inference NetworkPolicyHandlers"
41     cd controller/cmd && go build -o ../bin/main main.go
42
43 .PHONY: run-controller-cni-sec
44 run-controller-cni-sec:
45     @echo "Running the controller UNIX stream Inference NetworkPolicyHandlers"
46     cd controller/bin && ./main
47
48 .PHONY: build-controller-image
49 build-controller-image:
50     @echo "Building the controller docker image"
51     cd controller && docker build -t $(CONTROLLER_IMAGE_NAME) .
52
53 .PHONY: run-controller-image
54 run-controller-image:
55     @echo "Running the controller"
56     docker run --name controller -p $(CONTROLLER_PORT):9000 -d $(CONTROLLER_IMAGE_NAME):$(CONTROLLER_IMAGE_TAG)
57
58 .PHONY: stop-controller-image
59 stop-controller-image:
60     @echo "Stopping the controller"
61     docker kill controller
62
63 .PHONY: run-controller
64 run-controller:
65     @echo "Running the controller"
66     cd controller && java -jar bin/node-agent-controller-1.0-SNAPSHOT.jar
67
68 .PHONY: controller
69 controller:
70     @echo "Build and Run Controller"
```

```
1973 int classify(struct __sk_buff *skb){
1974     }else if (eth->h_proto == bpf_htons(ETH_P_IPV6)) {
1975         if (ipv6->nexthdr == IPPROTO_UDP) {
1976             // udp->dest == bpf_htons(LLMNR_EGRESS_LOCAL_MULTICAST)
1977         }
1978         if (actions.parse_dns_header_size(&cursor, false,
1979             return TC_DROP;
1980         void *dns_payload = cursor.data + sizeof(struct
1981         if ((void *) dns_payload + 1 > cursor.data_end)
1982         struct dns_header *dns = (struct dns_header *) (
1983         if (actions.parse_dns_payload_transport_udp(&cursor,
1984             return TC_DROP;
1985         }
1986         // reached app layer no offset processing required
1987         __u8 parse_flag = actions.parse_dns_payload_memor
1988         struct result_parse_dns_labels result = __parse_c
1989         // layer 7 rate limiting of the packet inside kernel
1990         __u16 dns_payload_size = udp_payload_exclude_header
1991         if (result.deep_scan_mirror) {
1992             #if DNS_RATE_LIMIT_VOLUME
1993                 __u8 dns_rate_limit_action = __dns_rate_l
1994                 // __u8 dns_rate_limit_action = 1;
1995                 if (dns_rate_limit_action == 0) return TC
1996             #endif
1997             #if DNS_RATE_LIMIT_TOKEN_BUCKET
1998                 if (__dns_rate_limit_tbi(&cursor, skb) ==
1999                     return TC_DROP;
2000             #endif
2001         }
2002         __u32 out = skb->iindex;
```

# Success Measure

- Response Speed
- Detection Accuracy
  - High Precision and Low False positives
- Volume of Data loss prior removal
- Scalability in distributed environments
- System Performance Impact
  - Kernel
  - Userspace



# Results and Evaluation

- Model Metrics
- Throughput comparisons (Active mode)
- Resources
  - Memory Usage
    - Security Agent memory usage at endpoints in data plane
  - CPU Flame Graph
    - eBPF Agent CPU Flame Graph

## Test Bench

CPU: Intel Xeon 6130

Memory: 8 GB

Linux Kernel: 6.12.4

Network Driver: netvsc

Bandwidth: 100 Gb/sec

Root QDISC: Fq\_Codel


Queues: 8 RX / TX

# DNN Model Metrics

Metric	Training	Validation
Accuracy	0.9973	0.9997
AUC	0.9997	0.9997
Loss	0.0099	0.0091
Precision	0.9959	0.9959
Recall	0.9987	0.9988


Table 5.1: Model Evaluation Metrics

Model Performance




Model Scores

# Throughput comparisons – Active Mode




Agent LRU Cache Read-Through Hit




ONNX Live Inferencing

# Throughput comparisons – Active Mode (continued)




Agent LRU Cache Read-Through Hit




ONNX Live Inferencing

# Resource Usage – eBPF Agent Flame Graph


## Userspace Busy-Polling overhead monitoring maps




## Kernel Epoll asynchronous I/O for maps monitoring



# Resource Usage - Memory



10,000 DNS Req / Sec



100,000 DNS Req / Sec



# Knowledge Gained

- Kernel Traffic control (QoS) Qdiscs (clsact, fq\_codel, codel, htb)
- Kernel TCP Congestion control (rene, cubic, BBR)
- Userspace-kernel synchronization (kernel spin locks, RCU, userspace mutex, atomic ref\_counters), kernel asynchronous I/O
- Kernel Security Layer (LSM, seccomp, TEE)
- Distributed Systems concepts intersection with system performance
  - Caching Write / Read-through policies
  - Caching Eviction Policies
  - Data Streaming
  - NUMA cache coherence → NetFlow Steering

# Future Work

- **Extend Support for DNS-over-TCP and Encrypted Tunnels:** Implement in-kernel eBPF-based detection for DNS-over-TCP replicating TCP state machine over kernel socket layer, paired with userspace DPI via Envoy proxy.
- **Add In-Kernel TLS Fingerprinting:** Use eBPF for TLS fingerprinting (e.g., JA3/JA4) to detect DNS exfiltration over TLS (DOH), DNS over mTLS, WireGuard.
- **Rate-Limiting Based on Volume and Throughput:** Integrate egress CSLACT-based dynamic rate limiting for DNS mass data breaches integrating EDT\_BPF, FQ\_CODEL and HTB QDISC's.
- **XDP-Based Flood Prevention:** Introduce XDP ingress filtering inside kernel to mitigate NXDOMAIN-based DNS water torture and DNS amplification attacks on the endpoint.

# References

- Singamaneni Krishnapriya and Sukhvinder Singh. "A comprehensive survey on advanced persistent threat (apt) detection techniques." *Computers, Materials and Continua*, 80(2):2675-2719, 2024.  
<https://www.sciencedirect.com/org/science/article/pii/S1546221824005952>
- Leyla Bilge, Engin Kirda, Christopher Kruegel, and Marco Balduzzi. "Exposure: Finding malicious domains using passive dns analysis." In *NDSS*, pages 1-17, 2011. [https://sites.cs.ucsb.edu/~chris/research/doc/ndss11\\_exposure.pdf](https://sites.cs.ucsb.edu/~chris/research/doc/ndss11_exposure.pdf)
- Manos Antonakakis, Roberto Perdisci, David Dagon, Wenke Lee, and Nick Feamster. "Building a dynamic reputation system for {DNS}." In *19th USENIX Security Symposium (USENIX Security 10)*, 2010.  
[https://www.usenix.org/legacy/event/sec10/tech/full\\_papers/Antonakakis.pdf](https://www.usenix.org/legacy/event/sec10/tech/full_papers/Antonakakis.pdf)
- Samaneh MahdaviFar, et al. "Lightweight hybrid detection of data exfiltration using dns based on machine learning." In *Proceedings of the 2021 11th International Conference on Communication and Network Security, ICCNS '21*, pages 80-86, 2022.  
<https://dl.acm.org/doi/10.1145/3507509.3507520>
- Jawad Ahmed, Hassan Habibi Gharakheili, Qasim Raza, Craig Russell, and Vijay Sivaraman. "Real-time detection of dns exfiltration and tunneling from enterprise networks." <https://ieeexplore.ieee.org/document/8717806>
- Jamal Hadi Salim. "Linux traffic control classifier-action subsystem architecture." *Proceedings of Netdev 0.1*, 2015.  
<https://people.netfilter.org/pablo/netdev0.1/papers/Linux-Traffic-Control-Classifer-Action-Subsystem-Architecture.pdf>
- Daniel Borkmann. "On getting tc classifier fully programmable with cls bpf." *Proceedings of netdev*, 1, 2016.  
<https://www.netdevconf.org/1.1/proceedings/papers/On-getting-tc-classifier-fully-programmable-with-cls-bpf.pdf>
- Kristijan Ziza, Pavle Vuletić, and Predrag Tadić. Dns exfiltration dataset, 2023. <https://data.mendeley.com/datasets/c4n7fckkz3/3>